



EÖTVÖS LORÁND UNIVERSITY

FACULTY OF INFORMATICS

DEPT. OF SOFTWARE TECHNOLOGY AND METHODOLOGY

Traffic Control and Infrastructure Organization Using Reinforcement Learning

Supervisor:

Dr. András Lőrincz
Head Senior Researcher

Author:

Dániel Kuknyó
IT for Autonomous Systems MsC.

Budapest, 2023

Contents

1	Introduction	3
1.1	Motivation	3
1.2	Goals and outline	4
1.3	Research hypotheses	5
2	Research Background	6
2.1	City Design	6
2.1.1	A brief history of the modern day street	6
2.1.2	The current state of city design	8
2.1.3	The stroad	9
2.2	Practical uses of the research	12
2.3	Deep learning in Infrastructure organization	14
2.3.1	The spectrum of expert areas	14
2.3.2	Infrastructure and building modeling	15
3	Methodology	17
3.1	Constructing a city	17
3.2	Intelligent driver model	19
3.3	Modes of operation	21
3.4	Reinforcement learning	22
3.4.1	Basics	22
3.4.2	Deep Q-Learning	24
4	Deep learning architecture	28
4.1	The modeling procedure	28
4.2	State space	30
4.2.1	From city to graph	30
4.2.2	Definition of the graph	31

4.3	Action space	32
4.4	Rewarding mechanism	34
4.4.1	Cost of infrastructure	35
4.4.2	Human factors	40
4.4.3	Motor vehicle aspect	41
4.4.4	Other considerations	41
4.5	Network architecture	42
4.5.1	Graph convolutional architecture	43
4.5.2	Embedding architecture	46
4.5.3	State-based architecture	47
5	Results	49
5.1	The testing environment	49
5.2	Episodic rewards	51
5.3	Rewards for nodes	55
5.4	Embeddings	57
5.5	Examples	60
6	Conclusion	65

Chapter 1

Introduction

1.1 Motivation

In recent years the traffic of cities became a rising topic with more and more city governments realizing that a motor focused city design is unsustainable. Cities all around Western Europe have started designing their cities around humans and public transit, and not around cars. E.g. Paris is planning to be a 15-minute city, Barcelona is incorporating the superblock design, Finland is creating non-intersecting paths between the common intersections and the Dutch are using intelligent traffic light systems to manipulate traffic flow in order to optimize it for both cars and pedestrians.

If one takes a look at how the Dutch infrastructure is designed, they will see that despite having less car lanes and overall less space for cars, the traffic flows more dynamically. This is thanks to the intelligent design of intersections, traffic lights and infrastructure. The methodology of this has been known ever since the 1970s, but the auto industry has been fighting against it ever since in order to gain more market. In Northern America one can observe what happens if a city is designed with cars in mind, requiring everyone to own a vehicle in order to participate in society. This results in worse accessibility to the public for the disabled, incapable, elderly and young people as well with less sense of community, relationships. The approach to creating walkable, human-centered and intelligent infrastructure that is optimal for both pedestrians and cars is laid out in detail in the literature such as Strong Towns, a non profit urban planning organization.

The aim for this research is to able to model a system of roads or city, and being

able to pinpoint mistakes made by development engineers, with the goal in mind to make the city more humanly livable, liquidate urban highways and make traffic infrastructure more efficient for the benefit of both humans and motorists. This is a key aspect in designing the future's cities and the way to have a more sustainable living ecosystem.

1.2 Goals and outline

The project will focus on building an environment that models traffic flow in a graph-based structure, then to train a reinforcement learning model to find the optimal configuration of the roads in order to transport the most cars in the most effective way possible. This is an expert area where the urban design principles play a key role: one can easily observe that the most effective way to transport as many cars as possible is if all roads are 8-lane highways. However it's also easy to see that it's miserable to live in a city where there are no quiet, auto-low streets and only 8-lane highways. This might be the best configuration for cars, but it would make the life of people living in the city absolutely horrible.

The reinforcement learning agent's goal is to learn how to create a city where it's nice to live for a human and fast to drive for a vehicle through building and destroying infrastructure. The goal of the rewarding system of the reinforcement learning environment is to express these ideas: building cost, traffic light/roundabout trade offs, how humans would feel living next to the road. The agent will have to decide where to build, destruct, or make roads 1-way to make the city's transportation flow dynamic but also make it livable for humans. The rewarding scheme is expected to reflect the principles laid down by Strong Towns and other urban planning organizations significant in the field like Happy Cities: Transforming Our Lives through urban design. The main metrics that are taken into account during the research are:

1. Time taken to arrive to the goal junction
2. Cost of infrastructure
3. Livability by humans (number of lanes, connection density)

1.3 Research hypotheses

If the research is successful, the following hypotheses can be verified or nullified:

1. Can traffic flow and driver behavior be modeled inside a reinforcement learning environment accurately enough for it to be representative of the real world?
2. Can predefined road configurations be optimized for traffic flow, human livability and cost efficiency at the same time?
3. Is there a single neural network architecture that can achieve the optimization in every configuration or there is a need for a more complex neural network as the complexity of the graph increases?

Chapter 2

Research Background

2.1 City Design

2.1.1 A brief history of the modern day street

The shape and fabric of cities is the result of hundreds or thousands of years of development, transformation and reshaping. Many geographical and historical events play a role in determining how a city's transportation infrastructure is designed. The industrial revolutions have been the key driving events of change in most countries as they have led to the transformation of the pedestrian centered traffic. The most significant factor in city design has always been some idea of happiness or philosophy. E.g. the *agora* in ancient Athens, which was essentially a main square. However the philosophy of good life was built into it as it can be observed through monuments, temples, the most important buildings and law courts that have been surrounding it. And this practice of building and designing architecture hasn't changed ever since [1]. Even today, when a new skyscraper is built, it will reflect the architect's, CEO's or other executive members' ideas of happiness or greatness. They will look at it, and decide if they like it: is it good to go for building or not. But what has changed since ancient Greece? What was the key motive that made us design cities focused around cars?

Philosophy has always been a pendulum oscillating between sky and earth - God and human. In the Middle ages, the focus was God. Then following it in the Renaissance, the human became the focus of inspiration. After that came the Baroque, again with God as the main motivator. After the Enlightenment and the birth of the modern day citizen, the first style trend was Rococo, focused around

humans. Around this time has religion lost its status as the main, unified cultural narrative as it can be observed in the works of Nietzsche and Derrida [2]. The last attempt at constituting God as the center of philosophy was in the style trend of Romantics. However as the modern era citizen of the Enlightenment wasn't religious anymore, the style attempted to find God in the human (*übermensch* - Nietzsche). This is expressed in various other fields, not just architecture or city design: art, poetry, statuary, astronomy and even mathematics. The loss of focus around God and religion in modern society can also be seen as a contributing factor to the shift in urban planning towards car centric cities. With the decline of traditional religious institutions and values, there has been a shift towards individualism and consumerism in Western societies. The car, as a symbol of personal freedom and mobility, has become a powerful cultural icon that reflects these values. In contrast, traditional urban planning, with its emphasis on public transportation, walkability, and community spaces, was often based on religious values and ideas about communal living. The decline of religious institutions and values also led to a decline in the importance of shared public spaces, such as parks, town squares, and community centers. These spaces, which were often designed to promote social cohesion and a sense of community, were gradually replaced by private spaces such as shopping malls and suburban subdivisions, which were designed around individual consumer preferences and the use of private cars.

Around a hundred years ago in the 1920s cars have begun to appear in the cities in greater and greater numbers. At first it was a new and expensive technology, hence it was mostly the wealthy class that owned them. At this point in time traffic rules were essentially nonexistent, so drivers could drive however they pleased. This quickly led to a large number of traffic related incidents. This was a massively bad PR for the auto industry, so they started influencing city design using their political and economical capital [3]. In school the children started to receive education about having to look both ways before crossing the road [4], because they are the ones that don't belong on the road and not the other way around. Peter Norton argues [5] that the crucial moment in the shift towards car centric urban design occurred in 1923, when Cincinnati residents were asked to vote on a proposal to impose a 25 mph speed limit on drivers. This measure would have made it difficult for car dealerships to sell cars, potentially leading to a decline in revenue, business, and political influence. As a result, representatives of the automobile industry launched

a concerted effort to defeat the measure, using tactics such as billboards, fliers, and even paying people to vote against it. Ultimately, the auto industry emerged victorious, gaining significant political and economic influence over the subsequent years.

2.1.2 The current state of city design

Ever since, car dependent infrastructure and lifestyle is thriving in the USA. Owning a private vehicle has become a minimum requirement to take part in society as the sprawl has led to a run down public transit and railway infrastructure. As of the current state, the automobile is dominating the transport in the USA and hence becoming one of the most important determinants of the American lifestyle [6]. This can be attributed to the way cities and residential buildings are designed. A typical American city has overwhelmingly two types of residential buildings: skyscrapers and high rise buildings downtown and low density suburban expansion around it. What's missing is a trade off between the two: mid density, mixed use housing that can at the same time house a sufficient amount of people in a relatively small area, but at the same time is relatively cheap to build in comparison with skyscrapers. This mid density residential housing planning would be the key guarantee to design walkable neighborhoods in the United States, however it's missing. Of course, there are some areas where this type of architecture can be found but this is not the norm, it's the exception. This is referred to as the missing middle problem [7]. This is partly because zoning laws in the USA don't make it possible to build this type of building. The remaining buildings are so spaced out because according US regulations there's a parking minimum for every building built. Parking minimums are requirements set by local governments or zoning codes that specify the minimum number of parking spaces that must be provided for a particular type of development. For example, a zoning code may require that a new apartment building must provide one parking space for every unit, or that a new office building must provide one parking space for every two employees.

From the way the cities are built comes car dependency as a consequence. However this has also led to public transit being underdeveloped: the housing density in the suburbs doesn't reach the threshold to make it worth making it being part of a bus or train route. No matter where the city puts the bus stop, the time

taken to get there on foot in the area is too much for people not living near it. Putting bus stops overall would require too much public transit infrastructure that would have to be subsidized from taxpayer money and would result in near empty buses. The optimal scenario is when a transit route can reach just enough people so that the maintenance and operation can be financed from ticket fares and minimal amount of subsidies. This is called the critical density. If there is a critical density reached a single transit stop can serve enough people, which will make it worth in fare revenue. It's also rare to find dedicated bus and cycling lanes in the USA.

Buses are generally stuck in traffic with the cars in rush hour, which makes it even less inviting for passengers. It's worthwhile to note here that adding lanes to an already existing road does not solve traffic, so it's not a solution to just make roads wide enough so that buses don't have to wait in line together with motor vehicles. As an example consider the recent expansion to the I-10 highway south of Houston known as the Katy Freeway. Together with a recent expansion, it's now at 26 lanes at its widest part at the time of writing but traffic jams are a regular occurrence. Adding lanes also reduces the safety of the road, as found by [8] who were analyzing the relationship between the road geometry and safety, and also [9], who found a positive correlation in the numbers of crashes and the number of lanes in urban highway road intersections. Dedicated bike lanes are also missing from the infrastructure: the most common cycling lanes one can find are painted bicycle gutters without any separation from high speed motor traffic that makes it very dangerous to ride a bicycle. This results in even more inequality in society.

Apart from the reasons mentioned before, car dependency has another different aspect, CO_2 emissions. A study has found that despite being 5% of the total population, the United States is responsible for 45% of all transportation CO_2 exhalation [10]. A different work in the field suggests that the most transport emissions could be reduced by converting low density housing to mid density [11]. Even more than converting mid density to high density.

2.1.3 The stroad

The stroad is a term created by a non-profit urban planning organization called Strong Towns that has been referred to and processed by several relevant pieces of literature like the book Confessions of a Recovering Engineer - Transportation for a

Strong Town [12]. The idea of the stroad lays down the guidelines which have to be reflected during the training phase of the agent as it highlights some key aspects of city design. To be able to understand the the stroad as a concept, it's necessary to first grasp the functions and purpose of a road and a street.

The road

A road is a means of fast transportation between different destinations, designed with wide and forgiving features to accommodate high-speed travel. In order to keep vehicles at a maximum speed for as long as possible, exits on the road are spaced far apart to reduce the number of intersections. This also helps to prevent drivers from needing to brake frequently, which can increase the risk of accidents. Turning lanes are often included on many roads to allow drivers to slow down before taking an exit, reducing the potential for danger. Adjacent to the road is a clear zone that is kept empty for emergency purposes, such as allowing vehicles to stop or for emergency vehicles to bypass a line of traffic. The font size on road signs is large, making them easy to read from a distance. The M3 highway in Hungary is an example of a road that has been optimized for safety and fast travel. As important pieces of infrastructure, roads play a crucial role in transportation and safety.

The street

The street is a densely populated environment that is primarily intended for human use. It serves as a complex space where city life occurs, and is therefore typically lively with people. The street is designed for low-speed travel, as it is considered a destination rather than a thoroughfare. High-speed traffic is incompatible with human activity, as it introduces an added element of danger to commuting. Buildings are often constructed directly adjacent to the sidewalk, making them easily accessible to pedestrians. The street itself can be likened to an outdoor room, where everything is scaled down to the size of human beings. It is designed to be welcoming and inviting to those who wish to walk, socialize, window shop, or simply breathe in some fresh air. Many shops, restaurants, cafes, and services have entrances and exits directly onto the street. Living near a well-designed street can also increase the value of a property. The difference in real estate prices between a flat near Váci Street versus one near the M3 highway is a testament to how much

society values human-centered design in urban planning. Streets play a vital role in tourism, livability, and overall quality of life within a city.

The stroad

The stroad is a name that's the result of the combination of the words street and road. It describes an urban motorway that is neither a street nor a road, despite trying to be both of them. The stroad is a street that's designed like a road, and doing so it fails to bring forth the advantages of either of them, but is struggling with the disadvantages of both of them. First the stroad is analyzed from the viewpoint of a road.

On the stroad there are many entrances and exits to businesses and services, like one would find on a street. However there are multiple lanes of high-speed motor traffic in both directions. This is an added factor of danger, as there are typically no dedicated turning lanes, so drivers have to pay constant attention if the car in front of them is taking a turn. Because of the high speeds involved this means that drivers have to come to a quick halt. It's not rare that stroads allow the 70 km/h speed limit, which drivers often surpass, because the stroad is straight and wide. The geometry of the street allows the drivers to reach higher speeds while still feeling safe, however this is a fake sense of safety, as there are more points of conflict than on a dedicated highway [13]. Roads can keep vehicles on high speeds because there are few exits. Traffic jams are also a common occurrence. Because of the large distances involved, walking feels very uncomfortable in this kind of environment. People don't feel safe walking next to high-speed motor traffic. It's also common that one can find large parking lots, which are also not inviting for the human as there's nothing to see or do there. It's also difficult to cross the stroad unless using a dedicated crosswalk and treading very carefully not to be involved in an accident. This is clearly a place for cars so it's not surprising that pretty much everyone drives here.

Next the stroad is evaluated from the point of view of the road. On the stroad the lanes are wide and there are a lot of them. Traffic flow is more chaotic than on the road because there are a lot of exits, hence cars change lanes more often, easily giving the change to dynomen situations: when there's no imminent danger, but the conditions can turn dangerous at any moment. This also reduces the speed of traffic flow. Cars changing lanes are the leading cause of the caterpillar or butterfly effect. Traffic volumes are also high as often there are no viable alternatives to driving: bus

and cycle lanes are nonexistent. This kind of car dependent infrastructure creates distances that only cars can bridge, but not everyone owns a car in society. This city design inherently widens the gap between the different layers of society.

Stroads are also very expensive, as they are built to a highway standard despite the fact that they can't be used as efficiently: with high speeds, while maintaining a strict level of safety as in the degree of slope, sharpness of turns and several other design decisions that have dedicated highway engineering and maintaining offices in almost any country. The density of intersections requires the city to install traffic lights, as they are the only type of intersection that can be operated with a related safety for a road of this scale. Traffic lights in the US can cost up to \$750.000 for an intersection. Roundabouts and right hand priority intersection are not viable options here because of the high speeds involved. Drivers need to see if they have to slow down ahead of time.

The majority of crashes within the United States happen on stroads. People walking and cycling are especially in danger as there's no infrastructure that separates them from high-speed motor traffic. The wide and open design of the stroad encourages drivers to drive fast and not pay attention to their surroundings [14]. As an example, during the coronavirus lockdowns in the beginning of 2020, people weren't traveling as much, hence miles traveled by drivers decreased. Despite 355 million less miles driven fatal car crashes per mile increased by up to 34% [15]. This leads to the conclusion that the bad design of the stroad isn't killing even more people is that it's because they are usually so jammed up that the drivers can't go fast enough to kill each other.

2.2 Practical uses of the research

Today a great percentage of a city's budget goes towards motor infrastructure planning, building, expanding and maintenance. Many cities mostly in North America keep building more new infrastructure that they can maintain causing a municipal insolvency. This is referred to as the "growth Ponzi scheme" named after the common investment fraud strategy in the literature [16]. It stems from the fact that maintenance costs start to appear after decades of road use, where it will be constant. It's easier to build roads than to keep fixing all the roads that are in bad condition. It also has political motivations as building a new road can serve as bet-

ter PR than fixing potholes. There's also a large rate of suburban sprawl in the US and Canada, and lower density results in more infrastructure being built. This is an urban design and an infrastructure design problem at the same time.

There are many solutions from the urban design side such as building more dense housing or replanning highway expansions to not destroy historically dense neighborhoods that serve a good amount of people with affordable housing and public transportation. It's also very difficult to run transit lines through the suburbs because of the low density. There's a lesser amount of people in the rational radius of the bus stop where it's worth taking the bus rather than to walk a long distance to a bus stop and then take the bus. For this reason it's not common for US and Canadian suburbs to have widely accessible transit lines going through them.

The other principle is infrastructure design, which is the main focus of this research. From an infrastructure design perspective careful preparation and planning must precede the execution of building a new piece of infrastructure or the expansion of an existing one. The city must thoroughly consider where to spend its finite funds. This is the aspect where this research can help in the decision making process. The developed software and framework considers a new or existing piece of infrastructure between junctions and optimizes it in terms of traffic flow and human-centered perspectives. The goal of the deep learning model is to find the optimal configuration which costs the least amount of money while providing a maximal amount of capacity expansion while still considering human factors in a traffic environment.

There can be three very distinct use cases of the piece of software that's being developed under the umbrella of the research:

1. Within-city: this use case comes down to optimizing an existing piece of infrastructure. This use case has some restrictions:
 - (a) No new roads can be built between junctions as this would necessarily mean destroying existing buildings.
 - (b) Only a limited number of lanes can be added to the existing roads.
 - (c) Junction types can be set freely.
2. Green field investment (Outside-city): in this problem the RL agent is given junctions and it's free to set everything between them. The hypothesis is that

there are no existing roads, and the junctions must be connected to each other in the best way possible:

- (a) Lanes and roads can be added freely.
 - (b) Maximum number of lanes can be anything (within rational limits).
 - (c) Junction types can be set freely.
3. Reparation mode (mixed): This is a mix of the previous two use cases. The RL agent is given a current traffic system and the task is to convert it to the best possible model while using a limited amount of money and resources.
- (a) The junctions are given.
 - (b) Lanes and roads can be added, existing ones can be destroyed.
 - (c) Junction types can be set freely.

If the human factors weren't taken into consideration the problem could be reduced into an optimization scheme where the target is to find the fastest traffic flow. The solution to this problem seems rather trivial: building 8-lane freeways between all the junctions. However a city like this can't be imagined within rational frames. The city budget and the experience of pedestrians must not be left out of the equation.

2.3 Deep learning in Infrastructure organization

2.3.1 The spectrum of expert areas

Predictions in new literature all point in one direction regarding traffic control and artificial intelligence: AI will most likely become the brain of infrastructure and traffic as well as many other expert areas. With the rise of autonomous vehicles this possibility seems closer than ever. There's a huge potential in operating, managing and monitoring transportation. The infrastructure is going to be very similar to that of the internet currently: a decentralized network that is operated by mesh and local servers, data flow being sensor readings coming to and going from the several vehicles on the road. There are currently existing studies on the viability of smart networks that infuse different kind of sensors with road infrastructure [17]. For public transportation the LOA and GOA frameworks lay down the principles of creating an ethical AI-driven solution for trams and trains respectively [18]. There are also

works that use traffic light governance systems enhanced by artificial intelligence to aid faster traffic flow [19]. Reinforcement learning has also been applied to control the cycle, phase and red/green ratio of traffic light systems [20]. All of the studies in the field can be placed on a spectrum that has the vehicle on one end and the road infrastructure on the other. Some works develop systems that enhance the vehicle, some are creating instruments for the road that directly influence traffic safety or speed like traffic lights. Others connect the road and the car together by infusing them with intelligent sensors that add a level of safety to the traveling experience.

2.3.2 Infrastructure and building modeling

Building information modeling (*BIM*) is the collective name for technologies and methods that aid the generation and modeling the digital representations of places in order to have a better overview of the building process. This introduces cost effectiveness, a level of safety and better ways to understand environmental impact. Today transportation infrastructure is a keystone of its economy by providing fast travel times for goods and orders, providing jobs thereby significantly helping economical and social development. In recent years a great need was shown for technologies that can increase the effectiveness of building motor transportation infrastructure. There was a significant development in the planning process for infrastructure. One of the results of this development is BIM. BIMs are essentially computer files that contain information on an object in order to help decision makers and planners with design, building and overview. Different kinds of dynamic data can also support BIMs like sensor measurements and signals from the building systems. Advances in fields like computer vision and artificial intelligence has given rise to a technological revolution in building information modeling. There have been attempts where building models were constructed using a combination of LIDAR and photometric sensor data. BIM data itself is not an AI-driven process but it can help aid AI-driven dynamic modeling by providing more accurate information on buildings and infrastructure that can be part of a data driven solution's environment model. BIM can be seen as both a 3D model and the information that it represents. In this case the model is more important as the representative of the information, not just solely a model file [21].

BIM by itself is not a very new technology. In fact it has been around since the 1980's. However combined with modern AI systems it can provide planners and

executors with more accurate information and in-depth modeling. However in the recent years the interest in construction tech has grown significantly. Construction tech is a larger collective term for technologies and methods that aim at automating construction and building processes. These can be mostly categorized into three categories [22]:

1. Planning: applications that aid the planning phase. These are mostly software solutions for the designing and simulation of an environment. Some companies are experimenting with AR (augmented reality) where the planned building and the current state of construction are merged into one simulated reality.
2. Sampling: these are software and hardware solutions that have the goal of extracting information from the site. They are mostly sampling and measuring devices that are now becoming far more advanced than a simple laser ruler. The LIDAR building scanners are an example. The output of these devices often get fed into the software that's responsible for planning and modeling.
3. Construction: robotic hardware tools that actively build or destroy infrastructure on the job site. Big leaps have been made in areas such as transportation of material, applying paint and cleaning surfaces.

According to some estimates, venture capital investment in construction technologies in the US reached as high as \$1 billion in 2018, which translates to a fourfold increase from 2013 [23]. This shows a large peak of interest in such technologies and gives reason for further research in such fields of expertise.

Chapter 3

Methodology

3.1 Constructing a city

Everything starts with a city that the agent will have to learn and optimize the structure of. At this point the city is best described by a graph. Graphs are high-level descriptors of connections between nodes. This gives an easy to use but highly customizable method to define a city to train on. To define a city using the graph there are two things needed to be defined:

1. List of nodes of the graph with coordinates: N_1, N_2, \dots, N_k .
2. List of connections between the nodes: $N_i \rightarrow N_j, i \in 1 \dots k, j \in 1 \dots k$.

These definitions together stored in a data file together are called a *construction protocol* from here on.

The developed software product provides an interface where the user can make a graph, describing the intersections (nodes) and roads (edges) of the city in question. This for example can be done in GeoGebra [24] and exported into a construction protocol in order to work as an input for the model. Below is an example simple city constructed:

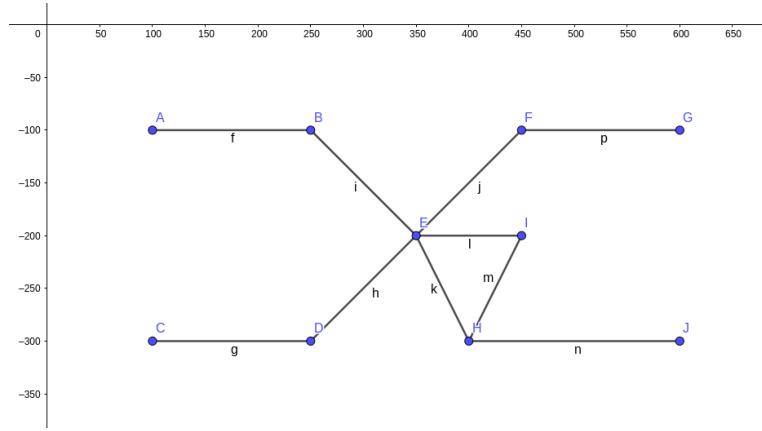


Figure 3.1: Graph-based representation of a city

The construction protocol is read into a graph, then the city is constructed from it using an inner representation. As a starting configuration every road is 2-way on the edges that have been set up in the graph. The vehicle rate and distribution can be controlled before starting the simulation. A constructed “starting” city according to the previously shown graph looks like as follows:

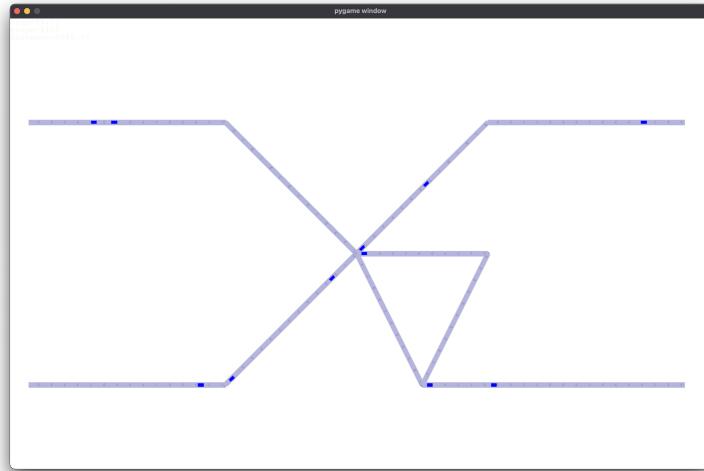


Figure 3.2: The simulation running the same city.

So far this is a simple setup for demonstration. The dark rectangles are vehicles on the road. The vehicles are passing from one entry point to another, without necessarily choosing the shortest path, or being evenly distributed among all the roads, just as one would find in real life. The driver model will incorporate an intelligent behavior, like slowing down after the car in front is slowing down or

gradually speeding up with a comfortable rate of acceleration after a light has turned green.

The road configuration is examined with multiple metrics like how many steps does it take for the roads to transport a given amount of cars or how much the road infrastructure would cost. If the agent is handed a road configuration it is expected to be able to find the optimal one, with the least cost, least unnecessary roads and fastest transportation for a given amount of cars. The cars are generated using a predefined distribution and each traffic junction has a separate entity of vehicle generator attached.

3.2 Intelligent driver model

The intelligent driver model (IDM) is a time-continuous car following model for the simulation of urban traffic. It describes the behavior of the drivers. and the positions of vehicles. The model incorporates the influence of one vehicle to another e.g. when some vehicle stops, the vehicle following it has to stop as well just like in real life. This property makes the model a good choice if the goal is to simulate real-world traffic as traffic jams, flash congestions and highway-like behavior is built into the model. The following 6 parameters influence how a driver behaves:

1. v_0 : the desired velocity of the vehicle
2. T : safe following time
3. a : maximum acceleration
4. b : comfortable deceleration
5. δ : acceleration exponent
6. s_0 : minimum distance between vehicles

The original model and the model that was used in this research was invented by Treiber, Hennecke and Helbing [25]. For the sake of simplicity the same parameters are used all over the simulation. The dynamics of the vehicle regarding acceleration and deceleration counting in the leading vehicle are then described by the following equations according to the authors:

$$\dot{x}_\alpha = \frac{\partial x_\alpha}{\partial t} = v_\alpha \quad (3.1)$$

$$\dot{v}_\alpha = a^{(\alpha)} \left[1 - \left(\frac{v_\alpha}{v_0^{(\alpha)}} \right)^\delta - \left(\frac{s^*(v_\alpha, \Delta v_\alpha)}{s_\alpha} \right)^2 \right] \quad (3.2)$$

Where x_α denotes the vehicle's position at time t for vehicle α . The deceleration depends on the ratio of the desired minimum gap and the actual gap:

$$s^*(v, \Delta v) = s_0^{(\alpha)} + s_1^\alpha \sqrt{\frac{v}{v_0^{(\alpha)}}} + T^\alpha v + \frac{v \Delta v}{2\sqrt{a^{(\alpha)} b^{(\alpha)}}} \quad (3.3)$$

Where T is the desired time headway. Using these interpolations of the variables and the positions of other vehicles it is possible at every time step to completely determine the position of every vehicle. Furthermore, vehicles can influence each others' positions, as there can be car pileups in an intersection with cars waiting for one another. This environment considers only cars of uniform size, without taking into account other types of vehicles like vans and semi trucks. Every driver is assumed to have the same skill set e.g. safe following time, comfortable deceleration and so on. This can be a fairly good approximation of a real world driver's parameters.

The simulation of the drivers' behavior is highly dependent on time, as they are regulated by differential equations that are implemented using computer code. Each time the agent will execute a training iteration, it will have to run a simulation. However, the simulation is not possible to run so many times in real-time. The game-time is implemented using discrete time steps, that can be sped up to arbitrary speed to run the simulation as fast as possible.

The IDM is very important in terms of the reinforcement learning problem, as traffic jams can be measured by calculating how long each car idles, how much distance did they cover while taking into account how many cars were spawned in the simulation. Part of the agent's reward is based on this term. If the cars covered more distance the reward will be higher, whereas if the cars idled for more time or traffic jams happened, the reward will be lower.

3.3 Modes of operation

There can be several different ways to use the framework depending on what the target of the prediction is. The modes depend on what actions are allowed. The more actions are allowed, the more the city will be redesigned from the ground up. The mode of operation is in a very strong connection with the use case of the framework (within-city, green-field or mixed) and essentially defines the types of actions that are allowed for the agent for specific usages.

1. Lane capacities: in this operation mode the nodes and junction types are fixed. The agent is only trying to predict how many lanes should be going from one intersection to another. The agent can add and remove lanes to existing infrastructure. This is the ideal mode of operation if there is a currently existing city, with already existing streets as it's not feasible to demolish residential apartment blocks just to build a new road between two intersections.
2. Lane capacities and roads: this is the same as the mode of operation mentioned before, the only difference is that the agent now has the capability to add a new road between junctions. This is the optimal configuration if it's the focused map segment is not already inside a city, but rather in a rural area where fast connection to nodes is a matter of hours of travel time on the highway.
3. Lane capacities and junctions: in this case the agent is not allowed to build any infrastructure that would mean destroying a piece of the city. It's only allowed to take the already existing roads and add or remove lanes from them, but not to build new roads where there isn't already one. The type of junction can be switched between right hand, roundabout and traffic light as it's still feasible to execute these types of modifications without needing to demolish residential properties.
4. Junctions only: the junctions only modality focuses on the nodes of the graph exclusively. The goal here is to determine what intersection types would be best fit to control traffic flow in an already determined road configuration. The agent has to choose between right handed, roundabout and traffic light for every junction in the map. Junctions only is the optimal setting in an environment where there's a narrow city e.g. like in Western Europe but there's

room for optimization in terms of intersections in order to determine the configuration that yields the fastest traffic flow.

5. Junctions and roads: this is the hardest task in terms of optimization and urban design. The mode allows the agent to set each intersection type, add lanes to already existing roads and add new roads between junctions. Newly built cities, already existing districts that are to be redesigned and between-city highways, country roads can also use this method of optimization. The research will focus on this mode as this is the most general.

3.4 Reinforcement learning

As this is a reinforcement learning based research so the following section will lay down the fundamental concepts of reinforcement learning. The definitions shown in this section are described by Sutton and Barto in their work Introduction to Reinforcement Learning [26].

3.4.1 Basics

Q-learning

The basic function for the model to estimate is the state-action value function, which translates formulates how profitable is it for the agent to be in state s , take action a and thereafter following policy π . The state-action value function the expected value of the sum of the cumulated discounted future reward, given state s_t and action a_t at time step t :

$$Q_\pi(s, a) = E_\pi [G_t \mid S_t = s, A_t = a] = E_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right] \quad (3.4)$$

Where γ is the discount factor, and R is the immediate reward received by the agent for its action. This expectation value is the one that the model will have to estimate. The update rule for the $Q(s, a)$ value is described in the Bellman optimality equation for the state-action value function. This is a very compact formula that uses the dynamics of the environment to calculate the $Q(s, a)$ value in a closed-form equation:

$$Q_\pi(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma V_\pi(s')] \quad (3.5)$$

Where the term $p(s', r | s, a)$ denotes the probability that the agent will receive reward r , end up in state s' , given that it's currently in state s and takes action a , and thereafter following the existing policy, π . $V_\pi(s')$ is the state-value function in the next state, s' which measures how good is it for the agent to be in state s and thereafter follow policy π . The issue with this formula is that despite it being very compact, it might take a lot of processing power to compute in complex environments, and also that the dynamics of the environment is not known for most cases (like in the case of traffic control). This is why it's reasonable to turn to deep learning methods to solve the task in order to estimate the $Q(s, a)$ value function. With the help of effective computer code and fast deep learning libraries, good accuracy and speed can be achieved.

Embeddings in deep learning

Embeddings are a common way to represent discrete categorical variables as continuous vectors. The most common use case of embeddings is in natural language processing, where words are represented as dense, high-dimensional vectors. The main principle behind learning embeddings is that similar features should have similar embedding representations. It is possible to train the embeddings' weights using backpropagation in a neural network. It is also common practice to use embeddings as inputs to other layers after they have been learned. The key feature of embeddings is that they are defined by their similarities and relationships. This comes in handy when performing a natural language processing task as the model can learn how different words are related to each other and thereby find context and meaning. Different words can have the same meaning despite having a very different character representation e.g. *fantastic* and *wonderful*. The neural network learns the similarity of these words by training their embeddings to become similar.

In this research the embeddings were used to represent the nodes of the city graph. In graph neural networks (GNN) embeddings are used to capture the structure and relationships within the graph, allowing the network to perform tasks such as link prediction and node classification. When a GNN learns an embedding it learns to map the nodes and edges into a continuous vector space such that similar

nodes and edges have similar representations and dissimilar nodes have dissimilar mappings. The embeddings in this case are able to capture topology, node attributes and the global structure as well. For example, GNNs can learn to represent nodes based on their degree, their position in the graph, the attributes of their neighbors, and the overall connectivity of the graph.

Graph convolutional networks (GCN)

Graph convolution is another approach to extract structural and representational information from a graph. Graph convolutional networks work very similarly as their image-based counterparts except that the input data is graph-structured. During a pass in a GCN layer, the information is not passed in a grid-like fashion as in traditional image-based convolution, rather information propagation works between neighboring nodes in the graph. GCNs can be used for node classification, link prediction, neighborhood detection, local and global pooling of nodes. The resulting information can be understood as embeddings of the nodes which capture the relationships and features of the nodes. Each node will have a custom embedding in the number of dimensions of the output layer's neurons.

The general propagation rule for a GCN is defined layer-wise where the input is a graph's adjacency matrix and the weights of the nodes in two separate data structures [27]:

$$H^{(l+1)} = \sigma \left(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)} \right) \quad (3.6)$$

Where $\sigma(\cdot)$ is the activation function, $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$ is the degree matrix of the graph, $\tilde{A} = A + I_N$ is the graph's adjacency matrix including self-connections, $W^{(l)}$ are the trainable weights of the layers and $H^{(l)}$ denotes the activations of the l -th layer in the network. The weights of the edges are passed as nonzero scalars in the correct entry in the adjacency matrix. This research uses GCN layers in order to extract information before dense fully connected layers.

3.4.2 Deep Q-Learning

Generic DQN

The deep Q-learning is an extension to the original method of calculating the state-action value function explicitly, using a (deep) neural network. In tasks where

the state is defined by an image of the gameplay screen, the go-to solution is a deep convolutional network that processes the raw pixel data and outputs the action to be taken by the agent. In the case of this research the environment passes the city's graph's adjacency matrix as the input to the Q_learning agent. In classical Q-learning the state-action values are updated iteratively by the following rule:

$$Q_{t+1}(s, a) \leftarrow Q_t(s, a) + \alpha \left(r(s, a) + \gamma \max_{a'} Q_t(s', a') - Q_t(s, a) \right) \quad (3.7)$$

Where α is the learning rate hyperparameter. and $r(s, a)$ is the immediate reward received for taking action a from state s . As stated before, this tabular approach is not possible as the environment has very large state space. However the Q-learning update rule can be directly implemented using a neural network. In this case instead of the Q-table, the model will be updating the weights of the neural network called parameters. The original temporal difference loss will be replaced by a loss function L :

$$L(\theta_t) = E_{(s,a) \sim D} \left[(y_t^Q - Q(s, a, \theta_t))^2 \right] \quad (3.8)$$

Where y_t^Q is the target value. The term D stands for an experience buffer which contains important data on the episode like state, action, next state, reward and dones (booleans to tell if the agent has terminated the episode with either win or fail). The experience buffer is sampled uniformly each time the model updates its weights and then emptied out. This originally was the idea of Long-Ji Jin [28] and is now widely accepted as a method of improving the model's performance. The loss function is calculated based on the sample and gives the experience for the agent to learn upon. This function measures how good the model performed on a given training iteration (episode). The procedure corresponding to the traditional Q-table update is gradient descent on the parameters of the neural network. It takes the following form:

$$\theta_{t+1} \leftarrow \theta_t + \alpha E_{(s,a) \sim D} \left[(y_t^Q - Q(s, a, \theta_t)) \nabla_\theta Q(s, a, \theta_t) \right] \quad (3.9)$$

Where ∇_θ is the gradient of the loss function.

Double DQN

In classical Q-learning and DQN, the algorithm always uses the same maximum values to select an action and to evaluate that action. This is what is referred to as optimistic estimation and this is something that is better to be avoided. The algorithm for Double Q-learning was originally proposed by Hasselt [29]. In this case two estimators are used synchronously to retrieve the predicted action value. The two estimators are referred to as target and local networks. When updating the state-action value function, one of the models is used to predict the best possible action, and the other one is used to evaluate the taken action. For this reason Double Q-learning has to store two Q-networks called Q^A and Q^B . The values from one network are copied over to the other one every defined couple of iterations. This results in lower variance and more robust models. The update rule is the same as in case of the DQN, but right now the model has to update the online network. However there's a slight change that has to be introduced as a result of having multiple estimators in the system. The target value in case of DDQN is:

$$Y_t^{DDQN} = r_{t+1} + \gamma Q(s_{t+1}, \text{argmax}_a Q(s_{t+1}, a, \theta_t), \theta^-) \quad (3.10)$$

Where θ^- is the set of parameters of the target network, the one for which the weights get updated every couple of iterations (defined by the programmer). The local network's weights get updated the same way as before. The update ratio is usually denoted by τ .

Dueling DQN

Before discussing the DQN viewpoint for dueling network architectures, it's important to define the advantage function. This measurement is one type of value function that's obtained by subtracting the state-value function from the state-action value function and represents how good is it to be in a state s , take action a and thereafter follow the given policy relative to the other actions that can be taken at that given state:

$$A(s, a) = Q(s, a) - V(s) \quad (3.11)$$

The dueling neural network architecture was first presented by Wang et. al. [30], which explicitly separates the inner calculation of the state-value function and the

state-action value function. The network's inputs are the same as in the cases before, the observations of the agent, and the hidden layers can be customized specific to the task as needed. However the main difference is in the output. The network will get two heads (output layers): one of the size 1, for estimating the state-value function of the current state and another one that's the same size as the action space (7 in case of traffic control) for estimating how good is it to take each possible action. Then after these values are calculated, the network reassembles them with the following forward mapping in order to maximize the Q-value:

$$Q(s, a, \theta, \alpha, \beta) = V(s, \theta, \beta) + A(s, a, \theta, \alpha) - \max_{a' \in |A|} A(s, a, \theta, \alpha) \quad (3.12)$$

The main principle of the dueling network architecture can be seen on the following graph:

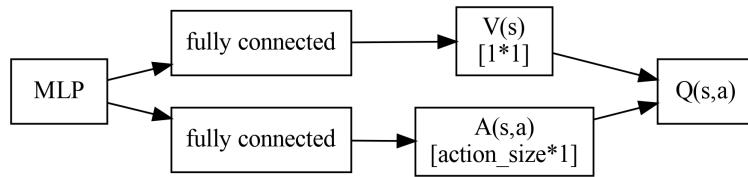


Figure 3.3: Dueling network architecture

The square node shape denotes deep learning layers with trainable weights and MLP stands for multilayer perceptron, a placeholder description of a neural network.

Chapter 4

Deep learning architecture

4.1 The modeling procedure

The modeling process describes the methodology to train the agent. Training translates to finding the weights of the neural network(s). The algorithm was chosen to be a very standard reinforcement learning framework with very little additions in order to restrict the agent and make sure that requirements are met.

The modeling structure is vastly dependent on the task that is given to it, therefore this section will only give a general overview of the inner workings of the deep learning modeling process. The main steps of a single iteration of simulation are given below:

Algorithm 1 The modeling process

Input:

1. A state vector x that describes connections between nodes in a graph G according to the current road configuration.
2. Number of episodes $n_episodes$.
3. Number of steps in episode n_steps .

Output:

1. The road configuration which is optimal for travel time, cost and livability.
2. The deep learning models with trained weights.

Process:

1. For $episode$ in $n_episodes$:

- (a) $A \leftarrow \emptyset$
- (b) $possible_nodes \leftarrow \{N_1, N_2, \dots, N_k\}$
- (c) Environment is reset

- (d) For $step$ in n_steps :

- i. Agent observes state vector x .
- ii. Environment chooses $start$ node $A \in possible_nodes$ if $A = \emptyset$.
- iii. Agent chooses end node $B \in possible_nodes$.
- iv. $possible_nodes \leftarrow possible_nodes \setminus A$.
- v. Agent chooses an $action a \in a_1, a_2, \dots, a_n$ from the set of valid actions depending on the $state$, $start$ and end .
- vi. The agent executes action a in the environment: $A \xrightarrow{a} B$.
- A. The environment builds or removes the piece of infrastructure.
- B. The simulation runs to a given stopping criterion (e.g. time elapsed, number of cars finished).
- C. The environment calculates and passes the agent the reward signal r and the modified state vector x .
- vii. The agent records $start$, end , $action$, $reward$ and executes a learning step based on the reward r if needed.
- viii. $possible_nodes \leftarrow \{N_1, N_2, \dots, N_k\}$ if $possible_nodes = \emptyset$:

It's easy to observe that the machine learning model has a particularly difficult job in this case. During the modeling process the agent has to step on the graph thereby generating a trajectory for itself. During each step it can choose the next node and what action to take. This is in order to encourage the agent for exploration and to make sure that it visits all the nodes in the graph. If the agent decides to build a lane an additional lane will be built from the current node to the next. If the agent changes the junction type the junction type will be changed in the target node.

Using the information extracted from the graph's structure around the starting node the agent predicts the ending node. Using the starting and ending node the agent will predict the action. The agent knows at each time step which are the possible actions with respect to the current state e.g. it can't remove a lane between two nodes where there are no lanes in existence currently or can't add a lane where the path has already reached maximum capacity.

The action it has to take in each iteration consists of multiple variables, not just a single one: $A \xrightarrow{a} B$. For this reason a highly specialized deep learning model configuration is needed. The input vector is the size of the state space as described by the node connections previously in this chapter.

4.2 State space

4.2.1 From city to graph

The state of the environment can be described in all time instants with a directed graph where the nodes of the graph are junctions and the edges are the roads connecting each junction. The weight of the edges refer to how many lanes are currently built between two junctions.

Initially, the focus will be on the representation of a specific kind of road that connects two nodes. The state between graph node A and B will have to be represented by a single number in every case. The bases that this scalar value will have to cover the number of lanes between $A \rightarrow B$ and the type of intersection in a node.

Second, the representation of the intersections: they are technically arbitrary but the possible range of values was chosen as $[1, 2, 3]$ from the cheapest to the most expensive: $\{1 : \text{righthand}, 2 : \text{roundabout}, 3 : \text{trafficlight}\}$.

The state of the intersection is in its essence a graph, so the following descriptions will have to reflect this principle. All descriptions will be given in terms of graph connections. The representation of the graph is the topic of a following section. For a state-vector element corresponding to the one-way connection between nodes A and B the possible values are as follows:

1. One-way, one-lane road between nodes A and B : $A \rightarrow B = 1$
2. One-way, two-lane road between nodes A and B : $A \rightarrow B = 2$
3. Two-way road, one lane in each direction between nodes A and B : $A \rightarrow B = 1$; $B \rightarrow A = 1$
4. Two way road, two lanes in each direction between nodes A and B : $A \rightarrow B = 2$; $B \rightarrow A = 2$
5. Right-hand intersection in node A : $A \rightarrow A = 1$
6. Roundabout in node A : $A \rightarrow A = 2$
7. Traffic light intersection in node A : $A \rightarrow A = 3$

The state-descriptor matrix keeps track of the connections in a directed fashion, storing the number of lanes from $A \rightarrow B$ and $B \rightarrow A$ in separate values (the graph is directed). A self-connection always refers to the type of junction that is present at the graph node. From now on the weight of a connection between nodes A and B in the state representation will be referred to as $x(A, B)$ or $x(N_1, N_2)$ e.g. what previously was mentioned as $A \rightarrow B = 1$ expressed in terms of the state descriptor matrix is the same as $x(A, B) = 1$.

4.2.2 Definition of the graph

A graph can be represented several different ways, each having its own advantages and disadvantages. For the reinforcement learning problem at hand the graph is represented using the graph's adjacency matrix. This opens up new possibilities and challenges as well. In the matrix-based approach each entry to the matrix describes a connection between two nodes, like a table. For n nodes the descriptor matrix x is of size $n * n$. The diagonal elements $x(N_i, N_i)$ describe the self-connections e.g. the type of intersection that can be found at junction N_i in the city. All other

connections $x(N_i, N_j)$, $i \neq j$ represent the number of lanes between nodes N_i, N_j . The representation is as follows:

	N_1	N_2	\dots	N_k
N_1	$x(N_1, N_1)$	$x(N_1, N_2)$	\dots	$x(N_1, N_k)$
N_2	$x(N_2, N_1)$	$x(N_2, N_2)$	\dots	$x(N_2, N_k)$
\vdots	\vdots	\vdots	\ddots	\vdots
N_k	$x(N_k, N_1)$	$x(N_k, N_2)$	\dots	$x(N_k, N_k)$

Table 4.1: Adjacency matrix representation

This approach yields advantages and disadvantages. The downside is that the matrix's size is fixed: there cannot be any edges that can be left out of the learning. If they are fixed on the value 0 to denote this, there is still a necessity to process them at each training iteration as deep learning libraries rarely have an option to manually skip a connection. This would lead to complicated program code and would raise more errors rather than gain execution speed. Also manually setting weights is not good practice in a deep learning focused approach.

The advantage is that with this method there is now a possibility to use two-dimensional or graph convolutional processing as the matrix can be processed as a single-channel monochrome image.

Another key insight is that the order of the columns doesn't actually matter: they can be ordered in any order so long until the connection descriptor elements are also in the right place. This can be a useful trait to exploit as the nodes can be added in order of closeness to a keypoint or sorted by any other comparison metric. The matrix would now become not just arbitrarily ordered, but also the node positions would also have a meaning to convey.

4.3 Action space

Depending on the starting configuration the agent will have a given set of junctions and some to no roads to start with. From here on the goal is to add/destroy roads, lanes and change intersection types in order to optimize the throughput, capacity and cost of the road and human-centeredness. Each action of the agent will take two graph nodes as a parameter and the roads will be configured accordingly.

Any action for node A and B will assume a single directed edge $A \rightarrow B$ starting from A and ending in B . There are cases where semantically it would make more sense to have more or less than 2 parameters but these cases can be generalized to an action with two parameters and hence be channeled into a neural network output of the same shape and size as all other cases.

The list of actions for the discrete action space:

1. $\text{add_lane}(A, B)$: Adds a single one way lane going from $A \rightarrow B$ in the state-descriptor: $x(A, B) += 1$. Only valid if $x(A, B) < \text{max_lanes}$.
2. $\text{remove_lane}(A, B)$: Removes a single lane going from $A \rightarrow B$ in the state-descriptor: $x(A, B) -= 1$. Only valid if $x(A, B) > 0$.
3. $\text{add_road}(A, B)$: Adds two lanes between the nodes A and B going $A \rightarrow B$ and $B \rightarrow A$. Only valid in case of nodes that don't have edges connecting them: $x(A, B) += 1$, $x(B, A) += 1$. Only valid if $x(A, B) < \text{max_lanes} \& x(B, A) < \text{max_lanes}$.
4. $\text{remove_road}(A, B)$: Removes two lanes between nodes A and B going from $A \rightarrow B$ and $B \rightarrow A$: $x(A, B) -= 1$, $x(B, A) += 1$. Only valid if $x(A, B) > 1 \& x(B, A) > 1$.
5. $\text{add_righthand}(A, A)$: Removes current traffic light or roundabout infrastructure to create a right-hand priority intersection in graph node A in the state-descriptor: $X(A, A) = 1$. Only valid if $X(A, A) \neq 1$.
6. $\text{add_roundabout}(A, A)$: Adds a roundabout to all roads entering the intersection of graph node A . Destroys current traffic light or right-hand intersection in the state-descriptor: $X(A, A) = 2$. Only valid if $X(A, A) \neq 2$.
7. $\text{add_trafficlight}(A, A)$: Creates a traffic light system to all roads entering the intersection of graph node A . Destroys current right-hand intersection or roundabout in the state-descriptor: $X(A, A) = 3$. Only valid if $X(A, A) \neq 3$.

The action space presented here corresponds to the operation mode '*Junctions and roads*'. The action space corresponding to '*Lane capacities only*' is $\{1, 2\}$. For '*Lane capacities and junctions*' the space is defined as $\{1, 2, 5, 6, 7\}$, for '*Lane capacities and roads*' the set of actions is $\{1, 2, 3, 4\}$ and for '*Junctions only*' it's $\{5, 6, 7\}$. It's worth

noting that in the practical implementation if the agent chooses to add a junction with two different nodes only node B will be considered. E.g. $\text{add_roundabout}(A, B)$ translates to $\text{add_roundabout}(B, B)$. This is a functional consideration: in case of junctions the action and the endpoint matters more than the graph connection.

4.4 Rewarding mechanism

At each time step in the process of reinforcement learning's agent-environment framework, the agent takes an action and as a result, the environment changes its state and returns a reward to the agent. The agent observes the environment through the state variable described in the previous section. This section will deal with the principles that will be reflected during sending the reward signal to the agent. The main directives like exact costs and traffic engineering perspectives are further explored in the literature review. It's worthy to note that all the reward variables are subject to manual fine-tuning and specific research so exact values are not given here.

The main goal for the agent is to build a city that is livable for humans by keeping the number of lanes, roads and infrastructure costs to the minimum at all times. If this constraint was not present, the agent could simply build 8-lane highways in streets that are a few hundred meters long. This would render being a pedestrian miserable. It's also necessary to not build roads between nodes that is not necessary from a transportation perspective. If there's a reasonable detour between two endpoints, the cars should take that road.

The reward is made up of differences in city assessment signals. The city assessment signal has three main components, each component is based on a different principle:

1. Cost of the infrastructure: building a long, 2-way road costs more than building a small one-way street between junctions close to each other, and this reward component reflects this. Building an intersection is more expensive than destroying it. This is a one time penalty for the reward that the agent will have to learn to accept in order to receive a long term bonus. Just as a city's government spends money on construction to further better the city's transportation on the long term.

2. Human factors: each reward incorporates a factor that reflects how livable a city is. If there are no traffic lights pedestrians can't cross the road. If roads are wide, long and tangled into one another the city will become hostile towards the residents. The reward component will be higher if the city is more livable, and lower if the city is less livable. This is the human score of the city.
3. Vehicular factors: it's a bad design principle to have too little capacity for the number of cars that are present in the city. The better the commuting experience for the vehicles, the higher the reward. This component awards faster travel times and dynamic flow and penalizes slow travel and traffic jams. This is called the vehicular score of the city.

Therefore the score for a city that assesses all these factors at time step t can be calculated as:

$$score_{city}^t = -cost_{infrastructure}^t + score_{human}^t + score_{vehicular}^t \quad (4.1)$$

The rewarding is handled by a separate module in the program as it's a large task to assess all the variables and configurations. During the first step in an episode a reward of 0 is given to the agent. The current wellness of the city gets saved into a variable and after the next step the agent will receive the difference of the city assessment score of the current step and the score of the previous step as the reward r_t :

$$r_t = score_{city}^{t-1} - score_{city}^t \quad (4.2)$$

Therefore the reward will be positive if the city score increases and negative if it decreases. If the agent can hold a steady positive reward signal the city will be scored higher and higher.

4.4.1 Cost of infrastructure

Inter-node infrastructure

The cost of building a lane has to be taken into consideration. It is more expensive to add new lanes to existing infrastructure than to build the first lanes at the beginning. This is because the sidewalk has to be destructed and then rebuilt. Also traffic flow is slower in the time of the construction because of closed roads, leading

to congestions. Out of the 7 possible actions of the agent 4 are in relation to building roads and lanes. The relationship between them is the following:

1. The cost of building and destroying is linearly dependent on the length of the segment. The length of a unit of lane is denoted l_{lane} .
2. The cost of building a unit (e.g. meter) of one-way lane is: c_{lane} .
3. The factor of how much an extra lane expansion costs is denoted λ where $\lambda > 1$.
4. The cost of building a lane factored in with the extra lane cost is taken as $b_{lane} = c_{lane} * l_{lane} * \lambda^{x(A,B)}$. So if there are no lanes between the two endpoints the cost is taken purely however after that there's an increasing penalty on additional lanes (this reflects the idea of expansions being more and more expensive as they get added later on in the development).
5. The cost of building a road is equivalent to building two lanes: $b_{road} = c_{lane} * l_{lane} * \lambda^{x(A,B)} + c_{lane} * l_{lane} * \lambda^{x(B,A)}$.
6. The cost of destroying a lane is less than building a lane. The coefficient parameter is denoted δ : $0 < \delta < 1$ Then the cost of destroying is: $d_{lane} = \delta b_{lane} = \delta * c_{lane} * l_{lane}$.
7. Then the cost of destroying a road: $d_{road} = 2\delta b_{lane}$.

The necessary inputs are: $c_{lane}, \delta, \lambda$.

Intra-node infrastructure

The next section will describe the cost of building different types of junctions. The types of junctions in increasing order of cost are:

1. Right-hand intersection: it is the cheapest, however it is also the most unsafe. Cheap to build and to maintain. Also inexpensive to convert to and from all other types of junctions.

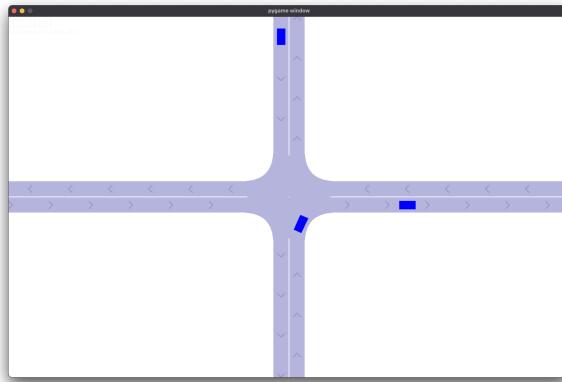


Figure 4.1: Right-hand intersection in the simulation window

2. Roundabout: a reasonable trade off between cost and safety, however it is difficult to destroy an intersection and build a roundabout in place of it as it requires widening the intersection. This is the best intersection in terms of dynamic traffic flow, but the worst in terms of human livability as it does not provide any safety features for pedestrians.

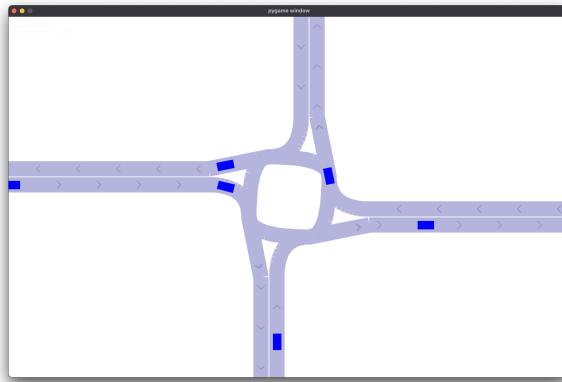


Figure 4.2: Roundabout in the simulation window

3. Traffic light: the most expensive and the safest type of junction is the traffic light as it requires dedicated infrastructure and the maintenance costs are higher than in case of all the other junctions. The traffic light is the best choice regarding human livability as it introduces an extra safety factor for humans. However traffic lights are known to cause traffic jams if the roads are not well planned. This type of intersection is difficult for the agent to figure out because as per environmental restrictions there has to be a certain number of inbound lanes to the junction in order to make it possible to build a traffic light intersection. The allowed number of inbound lanes in order to

build a traffic light intersection was [3,4,5,6]. The cycles of the traffic light are automatic and the agent has no influence on them. However this could make for the basis of a later extension of the research.

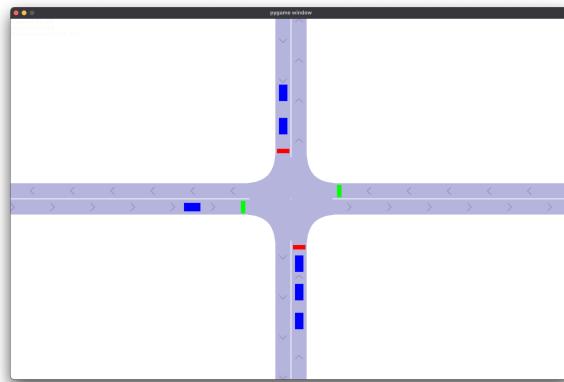


Figure 4.3: Traffic light intersection in the simulation window

As a graph node must be at all times assigned to a type of junction, and each type can be converted into any other type, the relationship between them can be defined by a triangle. The cost of converting from one to another can be parametrized and can be subject to change depending on how expensive building infrastructure is on a given terrain or economical environment. The conversion table is shown below:

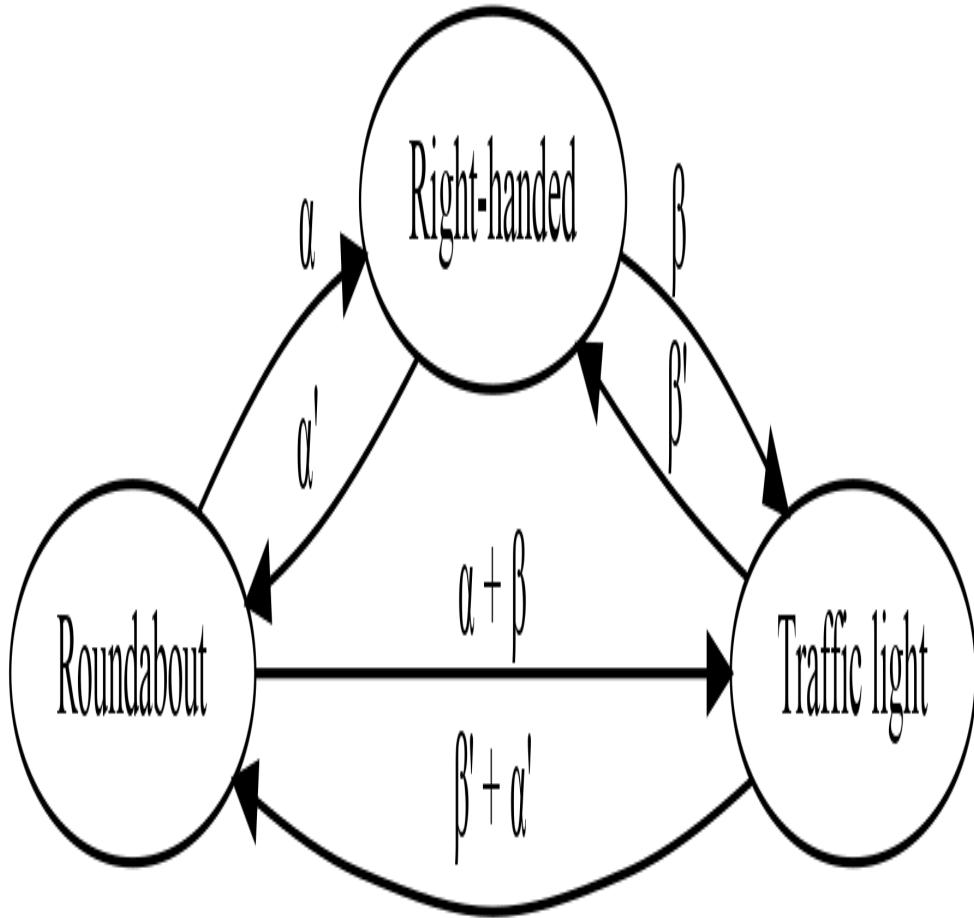


Figure 4.4: Conversion rates between junctions

The conversion will require the inputs to and from each type of intersection, as these are not necessarily functions on each other. Some dependency can be added later but that is independent of the current high-level model. The required inputs are: $\alpha, \alpha', \beta, \beta'$. Each variable represents a conversion cost between two types of infrastructure. The prime version of the variable is the reverse of the same building procedure. The prime ($'$) will always reflect a destruction e.g. demolishing an expensive piece of infrastructure and building a cheaper one in place. The non-prime variable denotes the building of a more expensive piece of infrastructure from a cheaper one.

By default conversions are defined between roundabouts and right-handed intersections, traffic lights and right-handed intersections. Converting a traffic light into a roundabout involves first removing the traffic lights, essentially converting the junction into a right-handed intersection and then rebuilding the right-handed

intersection into a roundabout. The same is true for converting a roundabout into a traffic light intersection: first the roundabout must be converted into a right-handed intersection and then traffic lights must be added to the intersection thereby finishing the conversion into a traffic light intersection. Therefore conversions between traffic lights and roundabouts are defined as $\alpha + \beta$ and $\alpha' + \beta'$.

There are limiting factors to different types of junctions e.g. how many roads it can intersect. The roundabout is famous for being able to intersect five or more roads without any problem because of the structure of it. Traffic lights can also have more roads than average coming in and out of it. For the most part right-hand intersections can handle at most four bidirectional roads. This is because human attention cannot be focused to so many places at the same time. More roads will translate to a more dangerous environment and will require highly dedicated infrastructure. It is possible, but it's rarely seen in the real world so it's not going to be considered in this research as it reflects bad design principles.

4.4.2 Human factors

Each training iteration the human factors get assessed according to the city map. Human factors considered in the project are:

1. How long roads are: roads too long get penalized. If the agent builds a road to the furthest part of the city from a given junction the face of the city will be too vastly modified. Each road longer than half of the spread of the city will result in a penalty.
2. How many lanes a road is: each additional lane is more expensive to build than the previous one. This is called expansion in construction terms. If the agent builds an additional lane above a given threshold a larger cost value will be accounted for.
3. Types of junctions: roundabouts are generally penalized and traffic lights are rewarded. Roundabouts are good for vehicles as they provide dynamic intersections but they don't provide possibilities for pedestrians to cross. The traffic light intersection is the exact opposite: it allows for pedestrians to cross while blocking high-speed motor traffic. Right-hand intersections get no penalties.

The human factor aspect represents how livable a city is for a human. This is a contradictory target with the motor vehicle aspect as motor infrastructure and human infrastructure are complementers of each other - the more motor vehicle infrastructure there is, the less human-centered is the architecture.

4.4.3 Motor vehicle aspect

Each training iteration the simulation is ran for a given amount of steps.

1. Car distance and vehicles spawned: through the simulation it's measured how much distance each car has taken and how long did they go. This value is then added to the reward, so a higher value will result in a larger reward.
2. No nodes alone: if the graph is continuous the agent receives a reward to learn to not leave out any junction from the traffic planning.
3. Nodes alone: if there's a node that has no connection to any other node, a penalty is added to the reward.

4.4.4 Other considerations

The agent can take a number of actions and some are incorrect in a given context. For this reason the agent's Q-values will get suppressed in the following cases:

1. Removing lane or road where there is currently nothing to remove (impossible).
2. Adding a lane or road where it has already reached maximum capacity (defined by n_lanes).
3. Adding a lane or road with the starting and ending nodes being the same (impossible).
4. Adding a junction where it is already built e.g. can't add a roundabout where it's already a roundabout (impossible).

This is an environmental restriction and is built into the framework to help the agent with learning the rules.

4.5 Network architecture

The input for the neural network is the graph of the city represented by an adjacency matrix. On each prediction pass the network has to predict three values: *end* and *action* (in this order). There were several implemented and tested architectures to achieve reward maximization such as dense networks, ensemble networks, embedding and graph convolutional networks. Below is the architectural description of the architectures that measured the best in the testing environment and served with enough insight so that comparison was worth it.

The architecture has to predict ending node B and action a at each iteration. The starting node is first given by the environment at random, then each iteration the previous ending node will be assigned as starting node. The naive implementation would be to use three separate neural networks predict each of the variables. This method was experimented on during training and the separate network architecture has been measured to perform poorly. In fact the single-net architectures' performance was so low that they didn't make it into the research paper. However they can still be tried and experimented on in the repository of the project. This architecture is erroneous in practice as two complete neural networks would have to be trained separately. The approach is redundant and can lead to information being lost and increased computation times.

In practice it's beneficial to make networks aware of each other's predictions. This approach is reasonable as there are dependencies between the variables: $A \Rightarrow B; A \Rightarrow a; B \Rightarrow a; x(A, B) \Rightarrow a$. The ending node depends on the starting node, the action depends on the starting starting node and ending node and the current state of the infrastructure that is currently built between them. If these dependencies are respected the true task can be learned better by the neural network.

The activation function used during the training was ReLu with an optimizer of Adadelta and learning rate of 0.001.

4.5.1 Graph convolutional architecture

End prediction

Below is the graph that is representative of the implemented deep learning model responsible for predicting the ending node based on structural information of the graph and the starting node that's given in the current learning iteration. Each square node represents a deep learning layer with trainable weights and each oval node represents some transformation in the neural network.

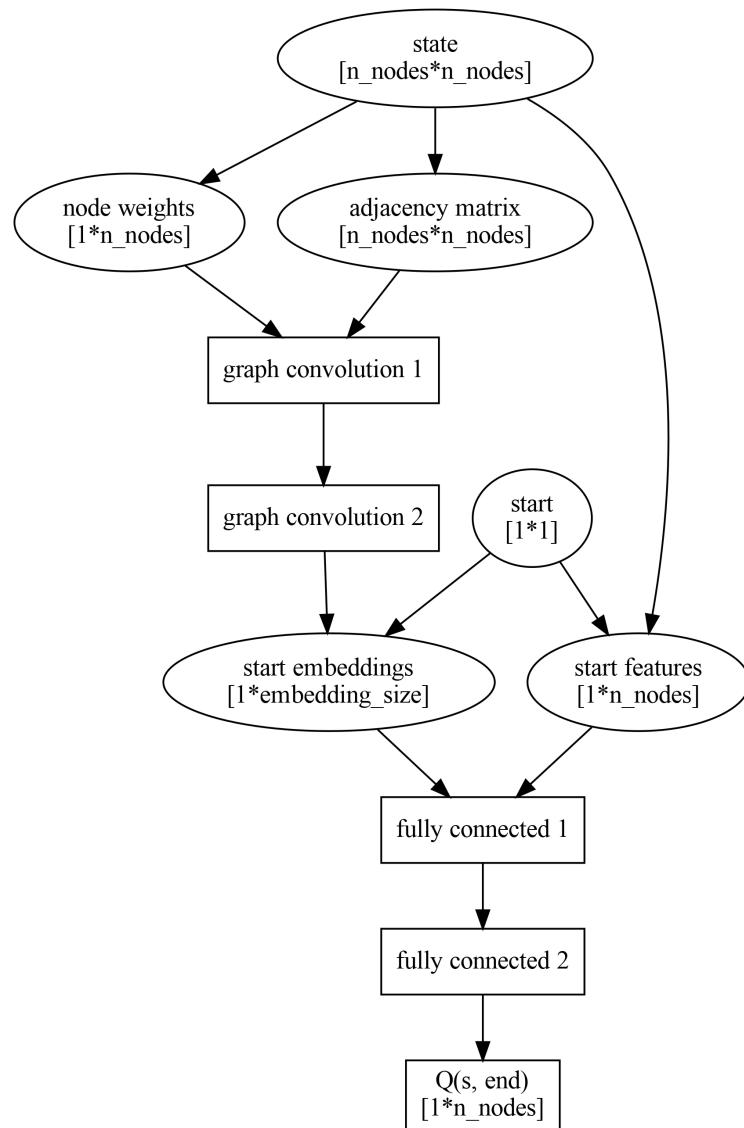


Figure 4.5: Graph convolutional network to predict the ending node

The input for the network is the adjacency matrix and the starting node. First the adjacency matrix gets separated into node weights, which is the state-definition matrix's main diagonal. The adjacency matrix is the state-definition matrix with

zeros on the main diagonal. The graph convolutional layers take these variables as inputs and produces an embedding for each node with a predefined size. The starting node's embedding and features from the state-definition matrix is selected using the starting node index. The embedding vector and the feature vector are concatenated and flattened and passed to densely connected layers. The output layer produces the state-action value function which represents how good is it to select an ending node from a given starting node. The agent will set the Q-values of the ending nodes that have already been visited to $-\infty$.

Action prediction

The second neural network is used to predict the action based on the state, start and end. For this reason it's slightly more complex but the general structure is the same as before. The two networks have separably trainable weights and the same number of layers. The number of neurons in the graph convolutional and fully connected layers can be customized as preferred. During most of the experiments each of the layers in the action prediction architecture had twice as many neurons as the corresponding layer in the end prediction architecture. The structure of the GCNN is as follows:

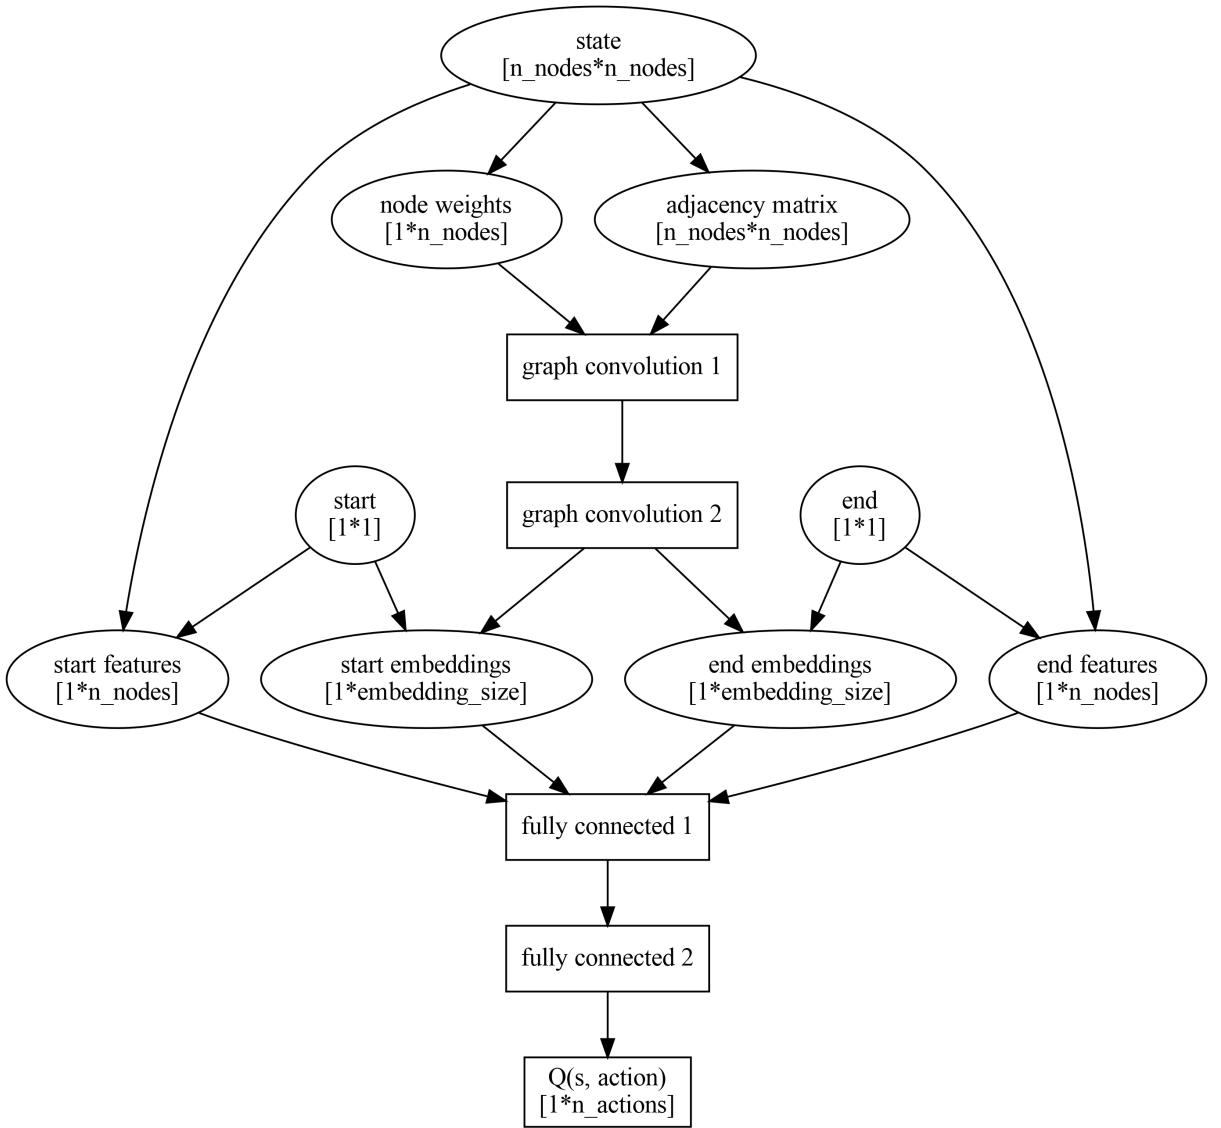


Figure 4.6: Graph convolutional network to predict the action

The inputs for the model are the state-definition matrix and the index of the starting and ending nodes. Just as before the state matrix gets separated into node weights and adjacency matrix in order to bring them to a shape that's appropriate for the graph convolutional layers. After the embeddings of the nodes have been extracted using the convolutional operations the correct embedding features are selected using the indices for start and end. The features of the start and end are selected from the matrix using the same indices. The four extracted vectors get concatenated into a single one that has the shape of $2*embedding_size+2*n_nodes$ and the resulting 1-dimensional vector gets passed to the fully connected layers. The amount of layers at each neural pass can be customized but during training two hidden layers were used. The output layer is used to predict the state-action value

function for the action which represents how lucrative is it for the agent to be in a certain state and given two nodes for the start and end take each action. The Q-values of the actions that cannot be taken for each start, end pair get suppressed to $-\infty$. This involves finding what the current junction is and how many lanes are currently built between the starting and ending node. If for example there are no lanes between the two junctions the agent cannot take the remove_lane action.

Some other considerations have also been taken into account while designing the neural networks for the modeling process like regularization through dropout which eliminates a node from the learning pass with a predefined probability. This forces the agent to generalize better. Another aspect of the models is that before each training iteration starts the model weights are initialized randomly using the Xavier method also known as Glorot initialization [31]. The resulting tensor will have weights uniformly sampled from $\mathcal{U}(-a, a)$ where

$$a = \text{gain} \sqrt{\frac{6}{\text{fan_in} + \text{fan_out}}} \quad (4.3)$$

Where gain is an optional scaling factor.

4.5.2 Embedding architecture

During running the experiments there were two other neural networks with a very similar architectures that created the basis for comparison. One that has used only embeddings (not graph convolutional) to make predictions, and one that used only the graph features from the state to determine the input for the dense connections. These models can be understood as submodels of the GCNN architecture with slight modifications so they will be introduced briefly.

The embedding architecture uses embeddings used in natural language processing to extract structural information from the graph. These embeddings serve as low dimensional representations of the graph's high dimensional and nonlinear structure. The embedding configuration to predict the ending node is as follows:

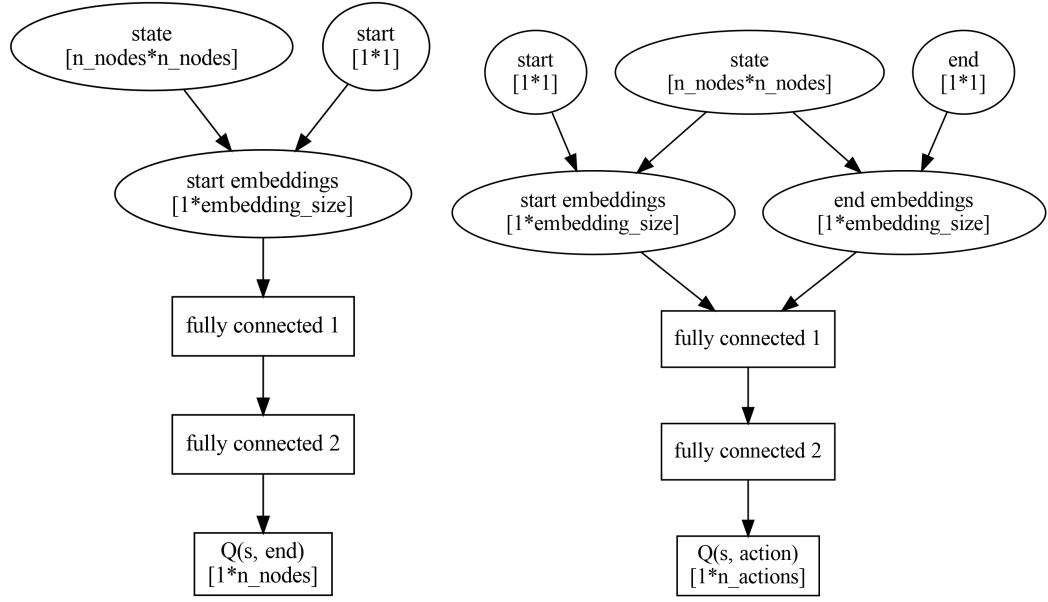


Figure 4.7: Embedding network to predict the ending node (left) and the action (right)

The input size for the densely connected layers is the same as the embedding size in case of the ending node prediction network and twice as much in the case of the action prediction network. Otherwise the network works the same way as its graph convolutional counterpart.

4.5.3 State-based architecture

The last of the configurations is the most simple one. This is a naive approach that served as an initial model and a baseline for comparison. The state-based architecture doesn't employ embeddings or convolution. It indexes into the state-definition matrix to acquire the features specific to a node: to which other nodes and how many connections it has and what type of intersection is currently set up in the junction that it represents. The state-definition based architecture to predict the starting and ending node:

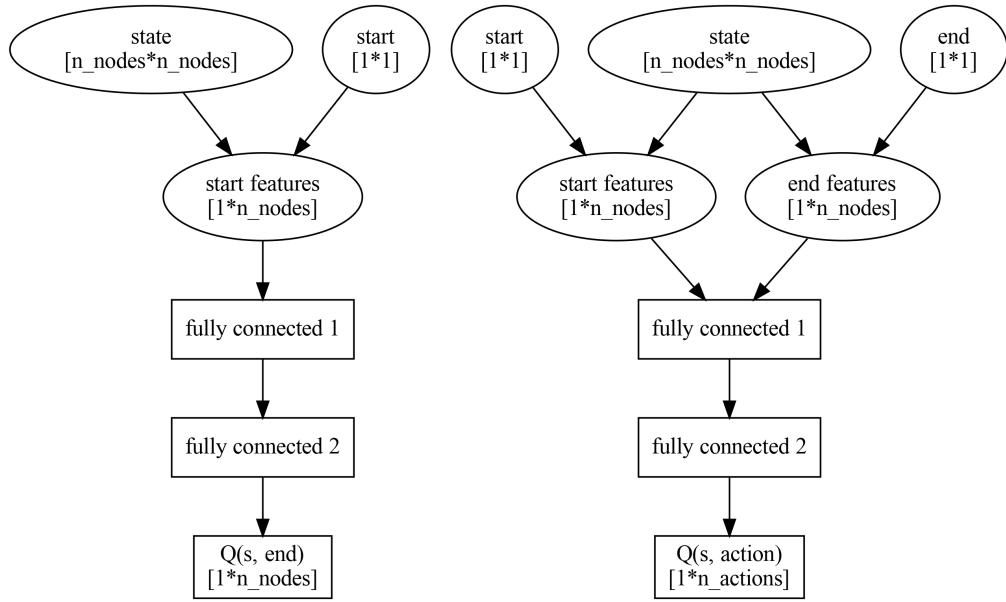


Figure 4.8: State-definition based architecture to predict the ending node (left) and the action (right)

Chapter 5

Results

The following section will deal with introducing the findings of the research. The machine learning task proved itself to be a very difficult one even for modern machine learning methods. Several of the most common approaches were implemented and tried in the environment however only a few of them proved successful. There were several factors that turned out to vastly influence the performance of the neural network such as costs, environmental restrictions and hard coded rules. The deep learning environment has more than 35 parameters, not counting the ones needed to define and train the neural network. The main metric that was used to measure the performance of the models was the cumulated reward during each learning epoch. The comparison of networks in this section will include networks trained with slightly different hyperparameters but with the same rewarding scheme for a fair criterion.

5.1 The testing environment

The main method for running the experiments was based on two notions:

1. The agent should be able to reconfigure an already suboptimal environment in order to make it more livable for humans while being able to provide faster transportation for cars. This corresponds to the idea of planning new expansions or destroying unnecessary roads in an already existing environment.
2. The agent should be able to create a road configuration from only the locations of the junctions, without having any roads or lanes added in the first place. This mode becomes useful in the planning phase of a green-field project

where there's a need to connect predefined places with the best possible cost effectiveness while providing fast travel times in a motor vehicle.

For the above reasons two testing environments were created. They are defined in such a way that there are multiple ways the agents can reach a fast transportation time, but creating a humanly livable environment at the same time is a difficult task. The number and locations of the nodes are exactly the same, the only difference is that there are some connections added to one of the environments. Out of the two distinct tasks the empty starting configuration can be classified as the harder one as the agent has no prior infrastructure to work with. The images of the starting configurations can be seen below:

1. The intersection that the agent will have to learn to optimize:

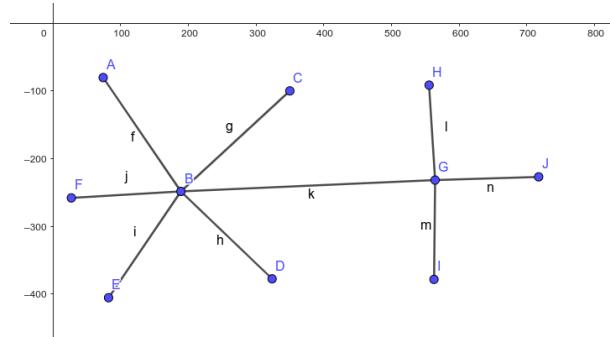


Figure 5.1: Predefined suboptimal environment: *test_5_intersection*

2. The intersection that the agent will have to learn to build from scratch:

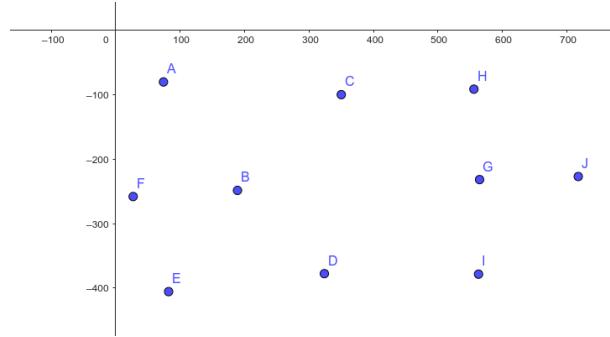


Figure 5.2: Predefined empty environment: *test_5_intersection_empty*

By default all roads have one lane going in each direction, and all junctions are configured to be right hand intersections. It's easy to see that in the first image the configuration is suboptimal as there are six roads meeting in one right-hand

junction. This is erroneous city design because of the danger of ambiguous right of way in the interchange. Another key thing to note is that this intersection is the bottleneck of the entire traffic system: by having six roads meet at the same place all of the vehicles will have to wait for their turn thereby likely causing a traffic jam. Traffic jams are measured and penalized by the environment.

Through the learning the agent is encouraged to either convert a right-hand intersection to a traffic light intersection for safety, or to convert it to a roundabout for speed. The traffic light intersection is more expensive to build however it adds an extra layer of safety to the configuration and for that it's rewarded. The more traffic light intersections does an environment have the more it scores in the human livability metric and the more roundabouts an environment has the more it scores on the vehicular aspect of the evaluation phase.

5.2 Episodic rewards

The following section will give an insight to the training rewards for each main type of agent employed in the learning. They have the same number of hidden layers, nodes, reward and environment configurations. The goal here is to decide which agent can capture the nuances of the traffic control problem better. The types of the agents are:

1. GCNN (Graph Convolutional Neural Network): a special architecture that uses embeddings from graph convolutional layers as well as node feature information indexed straight from the state-definition matrix.
2. ENN (Embedding Neural Network): a configuration which use solely node embeddings as the input for the densely connected layers. The embeddings here are the same as the ones used in natural language processing. The input for the embeddings is the state-definition matrix.
3. SNN: (State-based Neural Network): this one is the most naive configuration that uses only the adjacency matrix representation of the state to extract information on the nodes and connections of the graph.

During the training each agent was used to train for 1000 iterations on both of the environments that were described earlier. A batch size of 64 was used and the agents

updated their weights after every 8 iterations. For the end prediction networks the dense layers' number of neurons were [16,12,8] and for action prediction [32,16,8]. The GCN layers had 10 hidden units as well as the dimensions of the embeddings. As the agent will have to learn to optimize for long term rewards a discount factor of 0.99 was used all over the training. The strategy to optimize the rewards was epsilon greedy, where the agent is taking random actions with a given probability which is large at the beginning and decreases towards the end of the training. This kind of approach incentivizes exploration at start and exploitation in the end phase. Each training iteration the agent started out in node 0, and had 15 occasions to take some action. The largest possible reward was for adding a traffic light junction to one of the center nodes or any of the agent-built intersections with enough inbound roads. For each individual training the rewards of the agents were logged for each step and each action. The training for training in the suboptimal environment is the following:

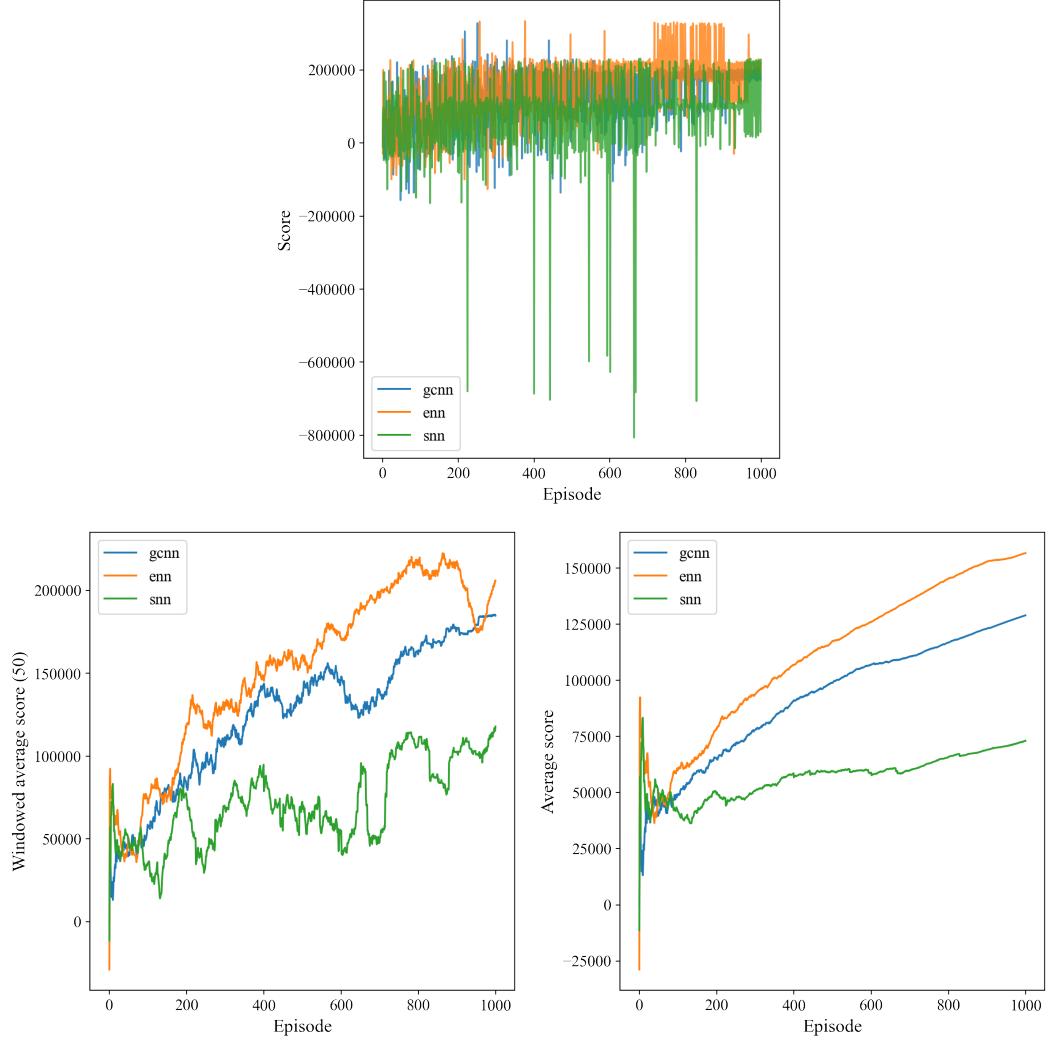


Figure 5.3: Training results in the *test_5_intersection* environment unsmoothed (top), smoothed with a 50 size window (bottom left) and total average (bottom right)

On the top is the reward for each episode. The rewards start out very noisy in the beginning but they stabilize around some point after a couple hundred iterations. In the bottom left there's the same exact data row smoothed with a window size of 50 and on the bottom right the same where for each episode the mean score was calculated. The mean reward \bar{r} at time step t : $\bar{r}_t = \frac{r_0 + r_1 + \dots + r_t}{t+1}$.

The worst performing model was the state-based architecture, where the embedding based and the graph convolutional models performed fairly well. The ENN model gained a lot of high scores during the late phase of the training giving it an advantage over the graph convolutional network. The largest negative scores could also be attributed to the SNN model, which also failed to converge around some reward value by the end of the 1000th iteration.

The other aspect to consider is the task of creating roads without any prior configuration given: the *test_5_intersection_empty* environment. The neural network agents were used to train with the same hyperparameters as before and the results were logged. The results are presented below:

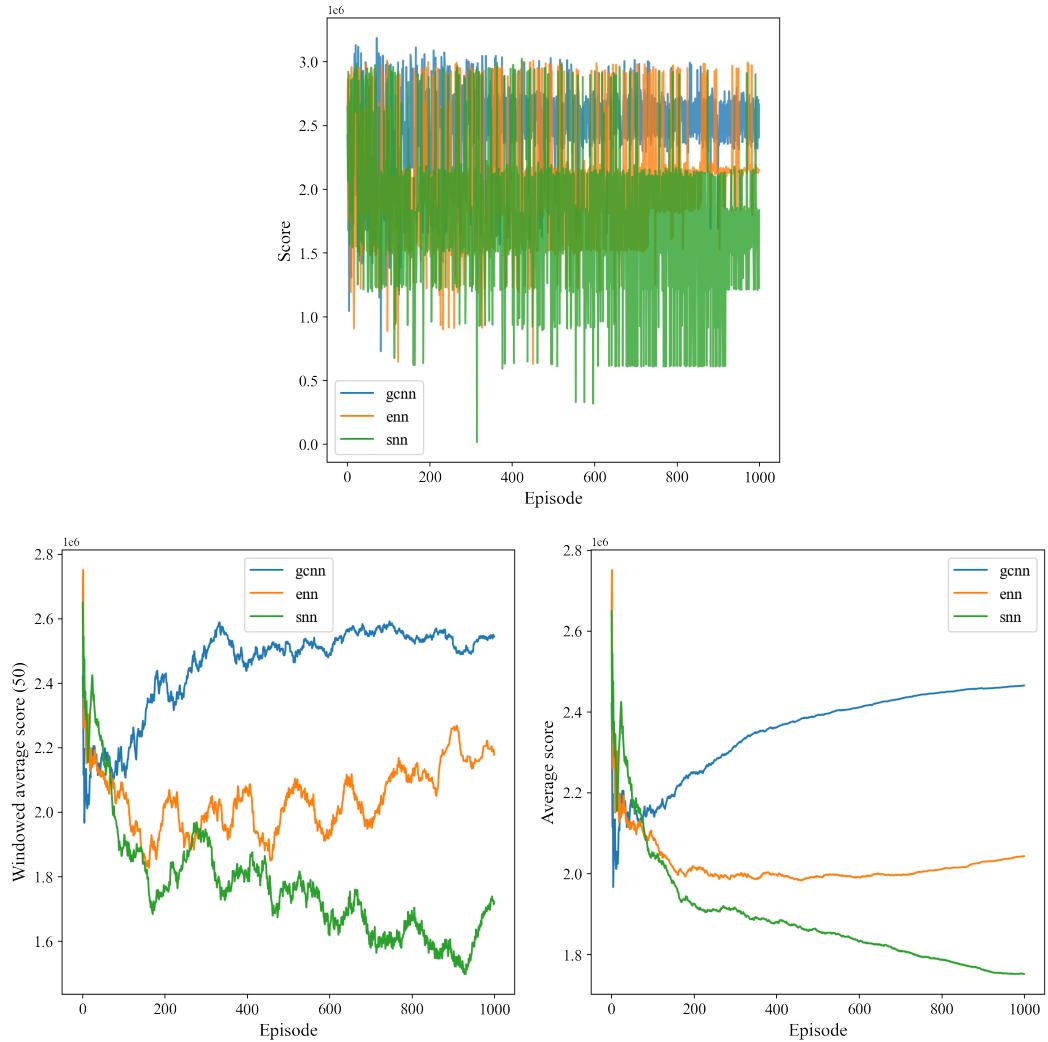


Figure 5.4: Training results in the *test_5_intersection* environment

This time the learning curves are diverging instead of showing a steadily increasing pace. The best performing model was the GCNN architecture which learned to increase rewards over time. The embedding based network was measured to be very unstable with the reward fluctuating between increasing and decreasing. Overall the reward has been slowly decreasing over time. Only the graph convolution based model learned how to increase rewards in the more difficult environment. Despite being noisier in terms of score deviation, still measured higher than the embedding

based network which has converged around a value that is still lower in score than the bottom mark of the GCNN model.

Another insight to note is that in this case it wasn't possible for the agents to add a traffic light intersection as the agent couldn't return to the same node twice.

5.3 Rewards for nodes

The next aspect of evaluation is how much reward each agent has received for building infrastructure to a given graph node. If the agent decided to take action a between the starting and ending nodes such that $A \xrightarrow{a} B$ the reward will be accredited to the cumulated reward of node B as this was the one that the agent has chosen (the starting node is chosen by the environment at first then the previous ending node becomes the starting node). This type of metric can give an insight to which nodes are more important in terms of infrastructure building and planning. This roughly translates to the idea of the state-value function $V(s)$ known in reinforcement learning which expresses how profitable is it for the agent to be in a certain state. The node based reward was calculated for each of the agents for both of the predefined environments. These results are from the same runs as seen in the previous chapter.

First the node based rewards for the *test_5_intersection* environment is presented for the three agents:

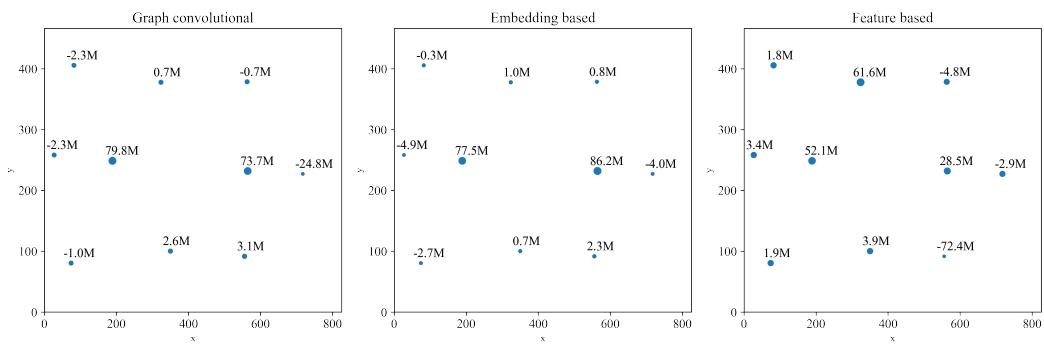


Figure 5.5: Cumulated reward for graph nodes in the *test_5_intersection* environment

On the above plot each dot represents a node in the graph with the actual (x, y) coordinates. The size of the node corresponds to how much reward did actions

directed towards that node receive during the learning. This information is also displayed next to the point in numeric form in the million scale.

Both the GCNN and the ENN models found the nodes in the middle of the other nodes to be more reward prone than the others on the outside. From the scoring we know that this is a better strategy than that of the feature based network. It's interesting to notice that some of the outer nodes have negative cumulative rewards for the GCNN and ENN models. The reason for this is that these are very likely to induce heavy infrastructure costs as the roads become longer. Even more so if the road is built from a distant node from the other side of the city. Another underlying factor is that the agent receives a large negative reward if a node is left out of the graph without any incoming or outgoing nodes. For the SNN the rewards are distributed more evenly among the inner and outer nodes which leads to the conclusion that the model didn't develop enough experience to distinguish the state-values of the nodes from each other.

The following plot is the same type of metric in the *test_5_intersection_empty* environment as seen in the plot below:

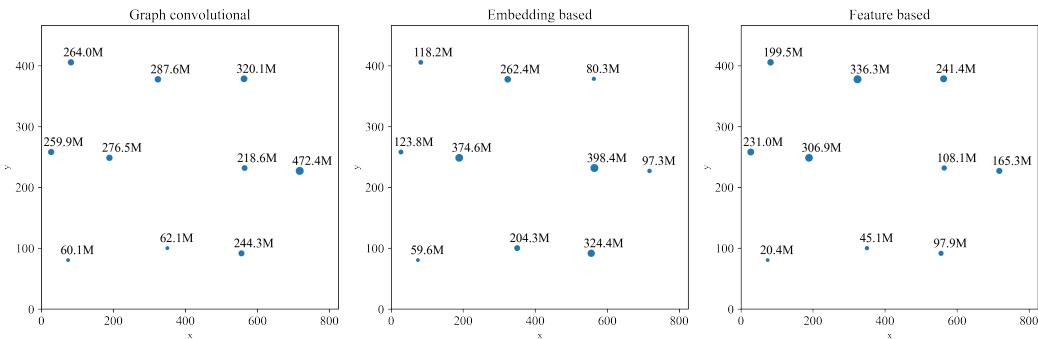


Figure 5.6: Cumulated reward for graph nodes in the *test_5_intersection* environment

This time the distributions are very uneven. In this testing environment the models had a more difficult job. Also it was more easy to gain reward as the initial reward is a highly negative value. For the GCNN model the most profitable node was the one on the very right which was unexpected given the nature of the task. In hindsight we know that this turned out to be a good strategy to optimize the rewards. In the case of the embedding based network the distribution of rewards is very similar to the one that was seen in the other environment with the exception of the bottom right node bringing an unusually high reward given the outer location.

For the state-based architecture the top left neighborhood of nodes yielded the highest reward. The bottom left and bottom middle nodes have been one of the least important ones in all the cases. This couldn't be attributed to them being far from the starting point as other nodes that are similarly far or even further have been able to receive an even higher amount of reward. What this means is that these nodes are not important from the aspect of traffic flow and human livability.

5.4 Embeddings

In this section the embeddings of the networks are visualized to compare their features in an untrained and trained state for the models where it is possible (GCNN, ENN). Embeddings capture structural information in a graph with the incentive that similar structures and neighborhoods should have similar embeddings.

An embedding size of 10 was used across the entire training procedure so the datasets are in a 10-dimensional space. To be able to visualize them a PCA based T-SNE model was applied using the same parameters across all the datasets containing the values of the embeddings. This method projects down a higher dimensional dataset into a lower dimensional space by aggregating axes into latent variables called components. To be able to visualize the dataset there has to be two components to be applied to the embeddings. There are in total two models with two neural networks each and two distinct states where the embeddings were captured.

1. GCNN, end prediction network:

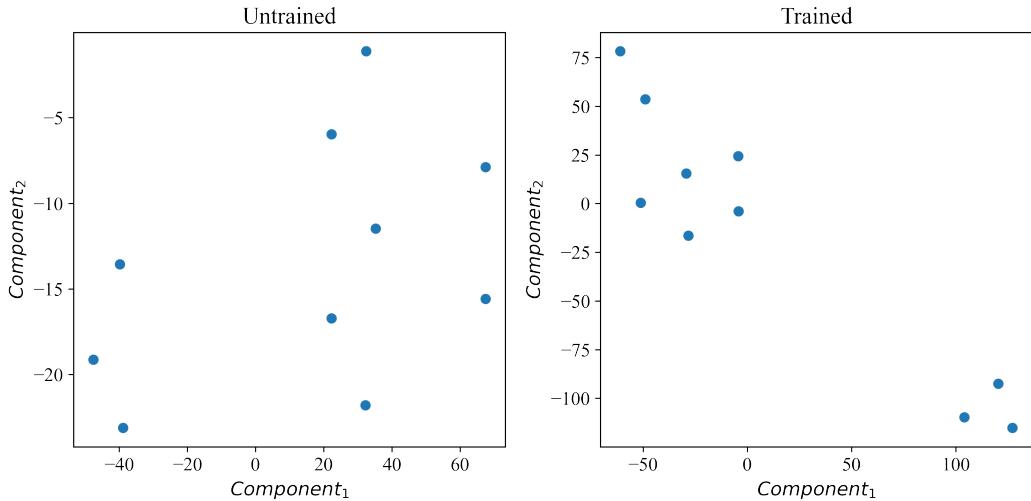


Figure 5.7: End embeddings for the GCNN model untrained (left) and trained (right)

2. GCNN, action prediction network:

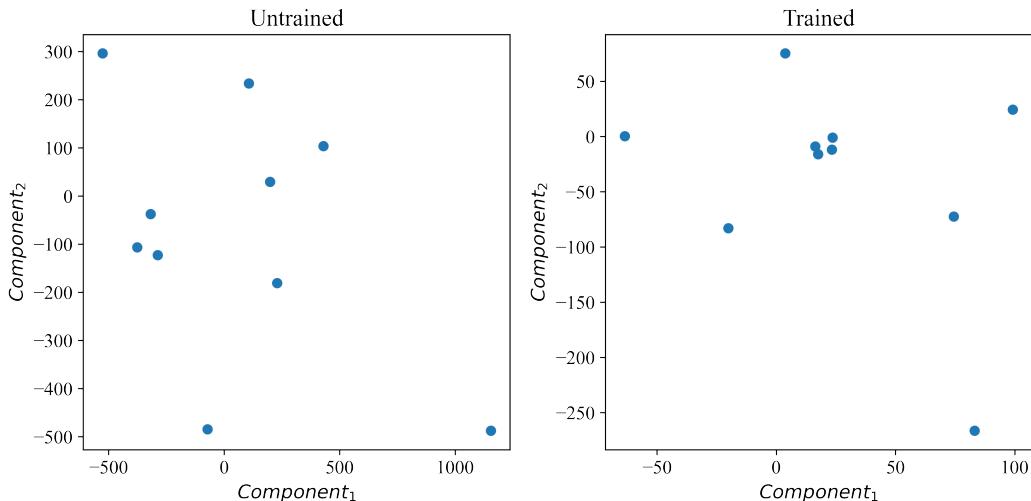


Figure 5.8: Action embeddings for the GCNN model untrained (left) and trained (right)

3. ENN, end prediction neural network:

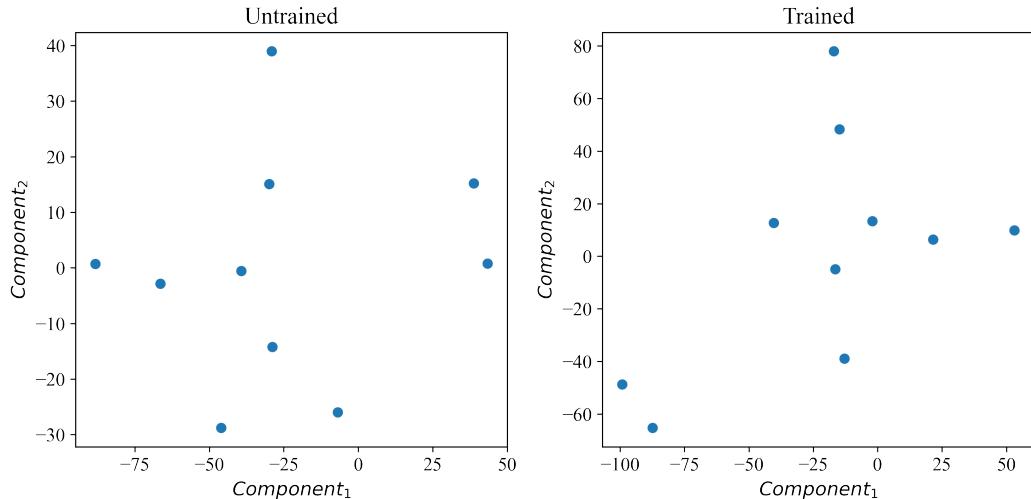


Figure 5.9: End embeddings for the ENN model untrained (left) and trained (right)

4. ENN, action prediction neural network:

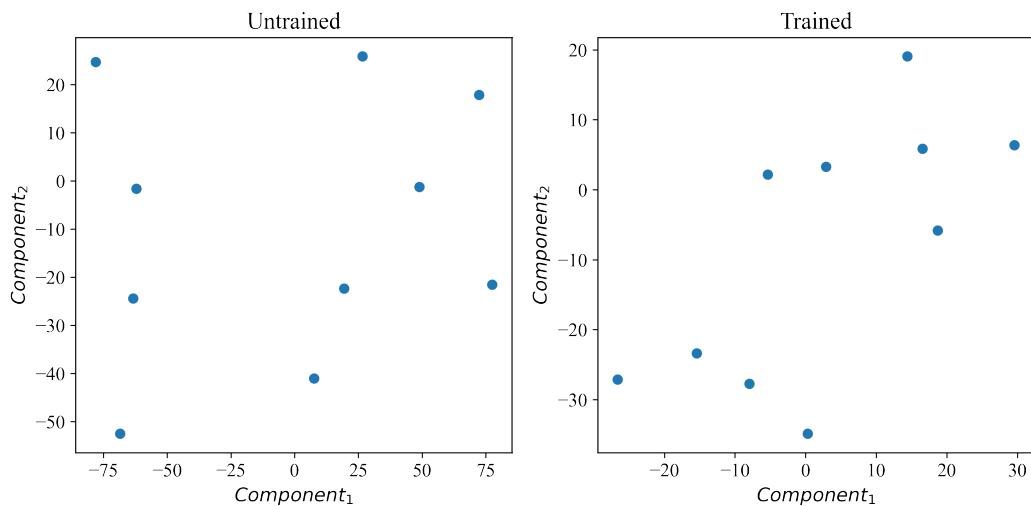


Figure 5.10: Action embeddings for the ENN model untrained (left) and trained (right)

It can be well noted that the neural networks seem to have some kind of idea on the similarity of the nodes in the untrained images. This is because the embedding layers' weights were initialized randomly and to get the embeddings the graph's adjacency matrix was passed as an input to the neural network. On these plots the closer the dots are the more similar they are according to the embedding layers. In the untrained images they are more spaced out and in their trained counterparts they are more closely bounded together. This is because the networks learned which nodes are similar to which other nodes in terms of structure and relationships to other nodes.

Of course structure is always changing, so always the beginning configuration of *test_5_intersection* was passed to the embedding layers to have an equal basis of comparison.

It's worth observing that the graph convolutional model made a larger change in the embeddings while training the weights. This leads to the conclusion that it has learned to distinguish them better and use the structural information from the graph in a more profitable manner. The fact that it has performed better in the harder task than the embedding based network proves this idea.

5.5 Examples

The following part is about giving examples on well scoring and badly designed cities. After the training process the agents were used to design a city for 15 discrete time steps, which means they had 15 opportunities to build or destroy infrastructure. One rule of building is that no new junction can be created, so when two roads cross the vehicles cannot change from one road to the other one in the intersection (essentially the crossing functions as a bridge). This was a way to simplify the learning task but can be the ground for the continuation of the research. If there are more lanes going between two junctions they show up as a single line in the simulation. The reason for this is that this research was focused on graphs and this is the most straightforward way to display a graph. Some examples with words of analysis are:

1. GCNN, *test_5_intersection*: this was one of the highest scoring cities of the entire training procedure. The agent has learned that if it connects the outer nodes then the traffic can be diverted between these nodes giving them a much larger vehicle score. It has also realized that in the center of the city many roads are meeting so it has added traffic light intersections to the central nodes thereby gaining a large positive reward in the human aspect scoring as well.

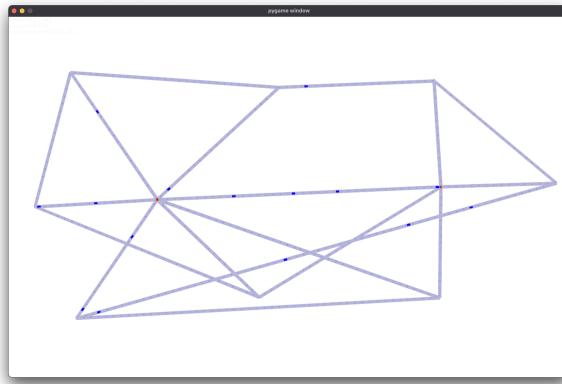


Figure 5.11: GCNN model, evaluation, *test_5_intersection*

2. GCNN, *test_5_intersection_empty*: this is a fairly well designed city for the harder task. The structure is again very similar to the one seen above: the outer nodes are connected for traffic flow, however the agent hasn't learned to add traffic light intersections to the inner junctions. For this goal it has to be highly adaptive as it only has 15 time steps to create a traffic light junction, and this type of intersection costs a lot and has the strictest criteria of them all.

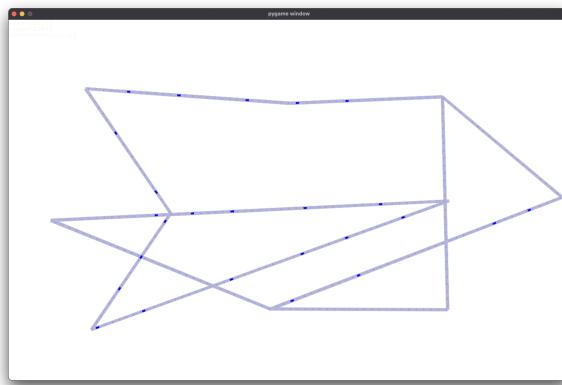


Figure 5.12: GCNN model, evaluation, *test_5_intersection_empty*

3. ENN, *test_5_intersection*: this is a purely traffic oriented map that is mainly focused around motor vehicles. The agent figured out that it gets a high reward if it makes way for high-speed traffic. There's an even higher reward if it also counts in the human aspect, however this was not the case in this experiment. The agent however learned that it can add roundabouts to the most used intersections thereby speeding up motor traffic even more. This takes a one time infrastructure building cost however speeds up the traffic from there on.

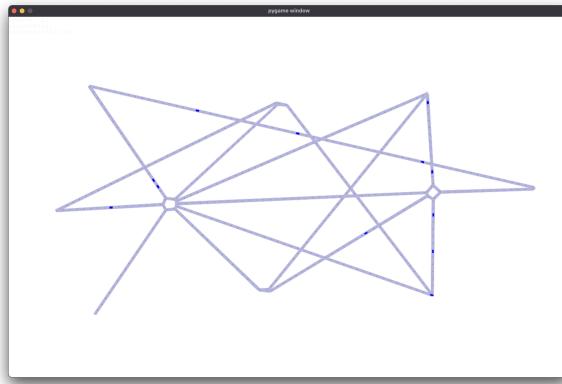


Figure 5.13: ENN model, evaluation, *test_5_intersection*

4. ENN, *test_5_intersection_empty*: as seen in the analysis of episodic rewards the ENN model had a clearly subpar performance on this testing environment. This can be noted in the organization of the map below as well: there are no special types of intersections, the graph isn't continuous and there are clearly entry points between which traveling includes needlessly long detours. This city scored very low during the testing and our human instincts support this evidence. This is not a city that anyone in their sane mind would design or live in.

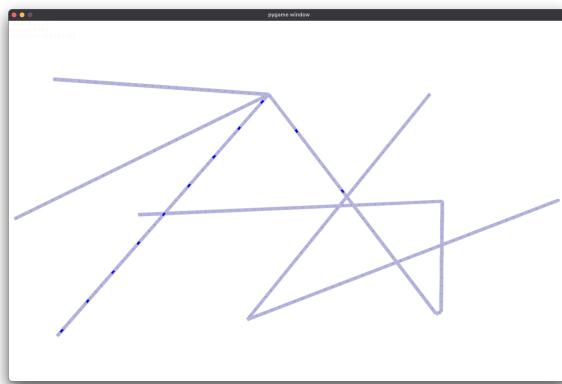


Figure 5.14: ENN model, evaluation, *test_5_intersection_empty*

5. SNN, *test_5_intersection*: this is a fairly well-designed city at first glance, however the agent has missed out on a lot of opportunities to gain a large reward so the others have outperformed it. It has added a lot of roundabouts to almost all the outer nodes which resulted in two things: a negative score as per the cost of building the infrastructure while not being able to significantly speed up traffic as these intersections have at most three inbound roads but

most often one or two. What the agent did right however is that it has added a roundabout to the intersection with the most inbound lanes in order to speed up traffic flow and a traffic light to the other central junction so that it gets the bonus for the human element of the scoring process. Despite the agent having the right idea, it has failed to produce good results on the long term.

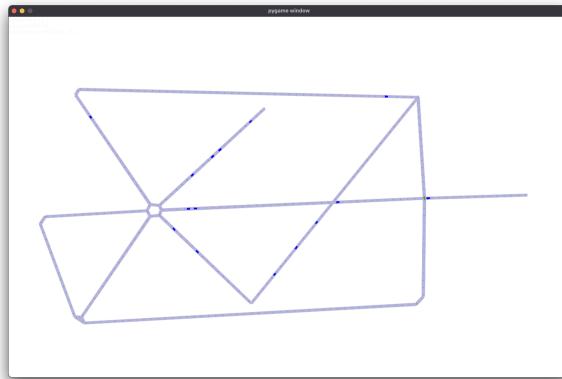


Figure 5.15: SNN model, evaluation, *test_5_intersection*

6. SNN, *test_5_intersection_empty*: one of the worst models in the history of the project was made on this test run. The agent had a steadily decreasing score during the training. It has failed to realize that adding roads will increase traffic speeds and connecting them in a way that they are reachable from one another will yield an even higher reward. It has failed to realize that adding roundabouts only makes sense if there are multiple inbound roads to the intersection so it has added three roundabouts with only two incoming lanes. The discontinuity of the graph and the fact that some nodes remained without any connections to other graphs also yielded a large negative reward. This is clearly a wrong configuration and is shown as an example to demonstrate the difference between an agent of superior and inferior performance in the testing environments.

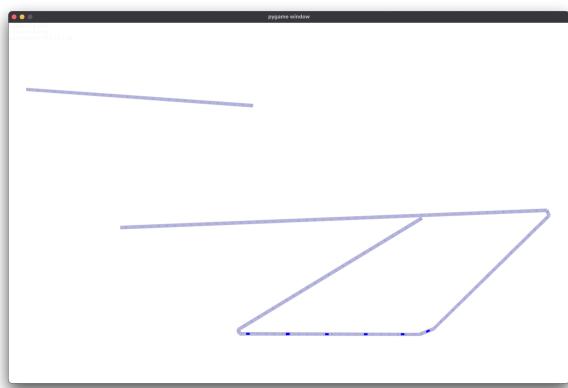


Figure 5.16: SNN model, evaluation, *test_5_intersection_empty*

Chapter 6

Conclusion

The task of traffic control using reinforcement learning has proved itself to be very difficult in terms of modeling, optimization and simulation. The need for a realistic simulation environment is large, however with a realistic simulation environment come a large number of variables which need to be handled by the machine learning model. It's a very time consuming problem to find the rewarding scheme that reflects what a human expects from a well-designed city while being able to combine it with the capability of a neural network to uncover insights from factors underlying the data. With the gained experience and understanding the following paragraphs will deal with answering the research hypotheses supposed in the beginning of the study.

Hypotheses

1. Can traffic flow and driver behavior be modeled inside a reinforcement learning environment accurately enough for it to be representative of the real world?
 - The bottleneck of the learning process is running the simulation after each action of the reinforcement learning agent. With the fairly simple simulation environment training times could take up to 24 hours to finish 1000 episodes. Even so sometimes the agents were measured to be underfit, showing a need for more actions inside one episode and for more episodes inside the training. To create an environment that accurately represents the real world and to play with it inside the environment would take heavy tolls on the time and budget of the laboratory, however technologically it is possible. The environment employed in this research wasn't

nearly as complicated as the real world, however it was enough for the agent to learn the main city design guidelines that it was designed to do. The hypothesis is supported.

2. Can predefined road configurations be optimized for traffic flow, human livability and cost efficiency at the same time?
 - It was shown in the research that if passed a configuration with infrastructure already present, the agent will learn to take the already existing structure and start building on top of it while using its advantages e.g. connecting the outer nodes in order to eliminate the middle intersections as the bottleneck of the process. Human livability factors were also learned through adding traffic light intersections and respecting the cost of adding new lanes or additional lanes next to existing ones. With the right reward configuration the human livability will be respected as much as needed for the current setup. The hypothesis is validated.
3. Is there a single neural network architecture that can achieve the optimization in every configuration or there is a need for a more complex neural network as the complexity of the graph increases?
 - There were experiments that were ran using a single, three headed neural network model. It didn't make it to the documentation because it was performing worse than random guessing. However it can still be downloaded in the repository of this project. The single network architectures didn't learn anything other than repeating the same exact action when prompted. One key advancement was separating the neural network into two distinct ones: one that predicts the ending node and another one that predicts the action. Random weight initialization and dropout have also helped dealing with their proneness to overfitting. Another key function was restricting the actions of the model to render invalid when a given action was not possible. All these evidences show that such a model has a lot of rules to learn and things such as exploration and restriction of actions helps it make the right decisions and improve the reward on the long run. The hypothesis is rejected.

Lastly

The purpose of this research was to develop a reinforcement learning model that can optimize traffic flow in a city by building roads, lanes and intersections between predefined nodes that are defined on a user friendly interface. The model will take into account how fast and dynamically does traffic flow and how welcoming a city is to its residents: the number of lanes and infrastructure cost have to be minimal while also minimizing the time taken to travel between junctions. To conduct the experiment, the machine learning model uses a simulated environment that incorporates an intelligent driver model in order to provide the most accurate mapping between real life and the modeled environment. The simulator proved itself to be sufficient for the small scale experiment however if the project is to be used in the industry there's a large need for an environment that can handle large amounts of data with a heavy graphics load.

Use cases of the research include optimizing current road structure and generating new paths between already existing cities or settlements that use the least amount of taxpayer money while also providing them with the fastest travel times. The developed framework can also be generalized for uses in public transport: using smart forecasting to optimize bus or train lanes to cover the most area, for as many people as possible. A future extension may include smart setting of traffic signal cycles to better divert traffic. Different types of vehicles and different driver behaviors would also help make the simulation environment more realistic in the sense that traffic flow would behave even more like in real life. It's also worth noting that this would add even more noise to an already very noisy system. One key thing that could prove itself useful is counting in the type of terrain that the roads are built on and being able to assign different infrastructure costs to different areas of land.

The most important insight was how different reinforcement learning algorithms can be made robust enough to handle a task with a lot of noise and variation. The methodology of creating such models is well known however this problem required a combination of regularization and restriction that is very task specific. The application of these methods was measured to yield results that serve with understanding that can be used in real-world projects as well. The application of this technology itself is very new with graph convolutional neural networks being only a few years old at the time of releasing this research. I sincerely hope that my research will serve

6. Conclusion

with the advancement of traffic control technologies for both the benefit of humans and motor vehicles.

List of Figures

3.1	Graph-based representation of a city	18
3.2	The simulation running the same city.	18
3.3	Dueling network architecture	27
4.1	Right-hand intersection in the simulation window	37
4.2	Roundabout in the simulation window	37
4.3	Traffic light intersection in the simulation window	38
4.4	Conversion rates between junctions	39
4.5	Graph convolutional network to predict the ending node	43
4.6	Graph convolutional network to predict the action	45
4.7	Embedding network to predict the ending node (left) and the action (right)	47
4.8	State-definition based architecture to predict the ending node (left) and the action (right)	48
5.1	Predefined suboptimal environment: <i>test_5_intersection</i>	50
5.2	Predefined empty environment: <i>test_5_intersection_empty</i>	50
5.3	Training results in the <i>test_5_intersection</i> environment unsmoothed (top), smoothed with a 50 size window (bottom left) and total average (bottom right)	53
5.4	Training results in the <i>test_5_intersection</i> environment	54
5.5	Cumulated reward for graph nodes in the <i>test_5_intersection</i> environment	55
5.6	Cumulated reward for graph nodes in the <i>test_5_intersection</i> environment	56
5.7	End embeddings for the GCNN model untrained (left) and trained (right)	58

5.8	Action embeddings for the GCNN model untrained (left) and trained (right)	58
5.9	End embeddings for the ENN model untrained (left) and trained (right)	59
5.10	Action embeddings for the ENN model untrained (left) and trained (right)	59
5.11	GCNN model, evaluation, <i>test_5_intersection</i>	61
5.12	GCNN model, evaluation, <i>test_5_intersection_empty</i>	61
5.13	ENN model, evaluation, <i>test_5_intersection</i>	62
5.14	ENN model, evaluation, <i>test_5_intersection_empty</i>	62
5.15	SNN model, evaluation, <i>test_5_intersection</i>	63
5.16	SNN model, evaluation, <i>test_5_intersection_empty</i>	64

List of Tables

4.1	Adjacency matrix representation	32
-----	---------------------------------	----